



Couchbase

# Architecture and Administration Basics

Data Modelling



1

# Data Access APIs



# Document Modeling Process

---

1. Understand your application requirements
2. Understand the Couchbase APIs
3. Map application access paths to best Couchbase API
4. Logical data modeling
5. Physical data modeling



## Understand your application requirements

---

- Inputs / Outputs
- Request rate, typical and peak
- Latency
- Consistency
- Scope of Aggregations
- Unique constraints
- Data set size & growth rate



# Couchbase APIs

Couchbase API	Latency	Throughput	Scalability	Applicability
Key/Value	500us–10ms	> 1M ops/sec	Highest	General, best for high throughput, highly latency sensitive workloads
N1QL	5ms+	> 40K qps	High	General, best for secondary lookups, pushing complex logic to database
Views	10–100ms	< 4K qps	Moderate	Aggregations, best for large scale aggregations (>1B docs) with low latency and moderate throughput requirements
Full Text Search	5ms+	> 20K qps	High	Text search, best for natural language queries, relevance ranking



# Mapping application requirements to APIs

Application Requirements	Couchbase API
<ul style="list-style-type: none"><li>• Very high throughput</li><li>• Latency sensitive</li><li>• Needs strong consistency</li><li>• Large data set / high growth</li></ul>	Key/Value
<ul style="list-style-type: none"><li>• Secondary key lookups</li><li>• Operational aggregations</li><li>• Filtered queries</li><li>• Ad-hoc queries</li></ul>	N1QL
<ul style="list-style-type: none"><li>• Report on well defined metrics</li><li>• Large scale aggregations</li><li>• Latency sensitive</li></ul>	Views
<ul style="list-style-type: none"><li>• Find patterns within text fields</li><li>• Provide search relevancy rankings</li></ul>	Full text search



2

## Data Modelling



# Goals of Data Modeling for N1QL

---

1. Define entities
2. Define relationships
3. Define document boundaries
  - Identifying parent and child objects
  - Deciding whether to embed child objects
4. Express relationships



# Defining Relationships

---

- Parent-child relationships
  - If we model the child as a separate document and not embedded, we have defined a relationship (parent-child)
- Independent relationships
  - Relationships between two independent objects
  - E.g. a person and a company where they work; deleting one does not delete the other (hopefully)



# Identifying Parent and Child Objects

---

- A parent object has an independent lifecycle
  - It is not deleted as part of deleting any other object
  - E.g. a registered user of a site
- A child object has a dependent lifecycle; It has no meaningful existence without its parent
  - It must be deleted when its parent is deleted
  - E.g. an invoice line item (child of the invoice object)
  - E.g. a comment on a blog (child of the blog object)



# Deciding Whether to Embed Child Objects

---

- Couchbase provides per-document atomicity
  - If the child and parent must be atomically updated or deleted together, embedding the child facilitates this
  - E.g. if an order line subtotal and order total must be updated together atomically, embedding the order line item facilitates this
- There is a performance tradeoff
  - Embedding the child makes it faster to read the parent together with all its children (single document fetch)
  - If the child has high cardinality (its parent has many instances of the child), embedding the child makes the parent bigger and slower to store and fetch
  - If the child has high cardinality (its parent has many instances of the child), embedding the child makes it more expensive to update the parent or the child



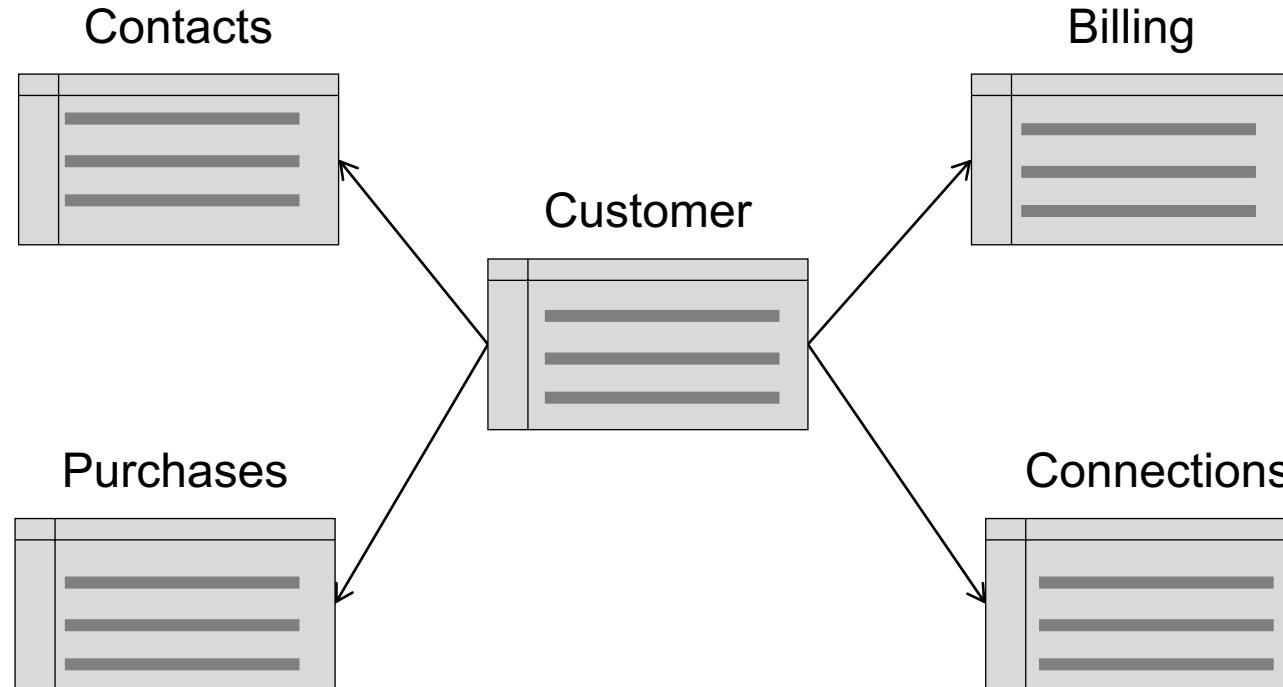
# Expressing Relationships

---

- 3 ways to express relationships in Couchbase
  1. Parent contains keys of children (outbound)
  2. Children contain key of parent (inbound)
  3. Both of the above (dual)
- High cardinality affects outbound relationships
  - Makes parent document bigger and slower
  - Makes it expensive to load a subset of relationships (e.g. paging through blog comments)



# Data Modeling with JSON Example





# Data Modeling with JSON Example

Table: Customer

CustomerID	Name	DOB
CBL2017	Jane Smith	1990-01-30

Customer DocumentKey: CBL2017

```
{  
  "Name" : "Jane Smith",  
  "DOB" : "1990-01-30"  
}
```

- The primary (CustomerID) becomes the DocumentKey
- Column name-Column value becomes KEY-VALUE pair.

OR

```
{  
  "Name" : {  
    "fname": "Jane",  
    "lname": "Smith"  
  },  
  "DOB" : "1990-01-30"  
}
```



# Data Modeling with JSON Example

Table: Customer

CustomerID	Name	DOB
CBL2017	Jane Smith	1990-01-30

Table: Billing

CustomerID	Type	Cardnum	Expiry
CBL2017	visa	5827...	2019-03

Customer DocumentKey: CBL2017

```
{  
  "Name" : "Jane Smith",  
  "DOB" : "1990-01-30",  
  "Billing" : [  
    {  
      "type" : "visa",  
      "cardnum" : "5827-2842-2847-3909",  
      "expiry" : "2019-03"  
    }  
  ]  
}
```

- Rich Structure & Relationships

- Billing information is stored as a sub-document
- There could be more than a single credit card. So, use an array.



# Data Modeling with JSON Example

Table: Customer

CustomerID	Name	DOB
CBL2017	Jane Smith	1990-01-30

Table: Billing

CustomerID	Type	Cardnum	Expiry
CBL2017	visa	5827...	2019-03
CBL2017	master	6274...	2018-12

Customer DocumentKey: CBL2017

```
{  
  "Name" : "Jane Smith",  
  "DOB"   : "1990-01-30",  
  "Billing" : [  
    {  
      "type"      : "visa",  
      "cardnum"  : "5827-2842-2847-3909",  
      "expiry"   : "2019-03"  
    },  
    {  
      "type"      : "master",  
      "cardnum"  : "6274-2542-5847-3949",  
      "expiry"   : "2018-12"  
    }  
  ]  
}
```

- **Value evolution**

- Simply add additional array element or update a value.



# Data Modeling with JSON Example

Table: Connections

CustomerID	ConnId	Name
CBL2017	XYZ987	Joe Smith
CBL2017	SKR007	Sam Smith

- **Structure evolution**

- Simply add new key-value pairs
- No downtime to add new KV pairs
- Applications can validate data
- Structure evolution over time.

- **Relations via Reference**

Customer Document Key: CBL2017

```
{  
  "Name" : "Jane Smith",  
  "DOB" : "1990-01-30",  
  "Billing" : [  
    {  
      "type" : "visa",  
      "cardnum" : "5827-2842-2847-3909",  
      "expiry" : "2019-03"  
    },  
    {  
      "type" : "master",  
      "cardnum" : "6274-2542-5847-3949",  
      "expiry" : "2018-12"  
    }  
  "Connections" : [  
    {  
      "ConnId" : "XYZ987",  
      "Name" : "Joe Smith"  
    },  
    {  
      "ConnId" : "SKR007",  
      "Name" : "Sam Smith"  
    }  
}
```



# Data Modeling with JSON Example

CustomerID	ConnId	Name
CBL2017	XYZ987	Joe Smith
CBL2017	SKR007	Sam Smith

CustomerID	Type	Cardnum	Expiry
CBL2017	visa	5827...	2019-03
CBL2017	master	6274...	2018-12

Contacts

Billing

Customer

CustomerID	Name	DOB
CBL2017	Jane Smith	1990-01-30

Purchases

CustomerID	item	amt
CBL2017	mac	2823.52
CBL2017	ipad2	623.52

CustomerID	ConnId	Name
CBL2017	XYZ987	Joe Smith
CBL2017	SKR007	Sam Smith

```
{
  "Name" : "Jane Smith",
  "DOB" : "1990-01-30",
  "Billing" : [
    {
      "type" : "visa",
      "cardnum" : "5827-2842-2847-3909",
      "expiry" : "2019-03"
    },
    {
      "type" : "master",
      "cardnum" : "6274-2842-2847-3909",
      "expiry" : "2019-03"
    }
  ],
  "Connections" : [
    {
      "CustId" : "XYZ987",
      "Name" : "Joe Smith"
    },
    {
      "CustId" : "PQR823",
      "Name" : "Dylan Smith"
    },
    {
      "CustId" : "PQR823",
      "Name" : "Dylan Smith"
    }
  ],
  "Purchases" : [
    { "id":12, item: "mac", "amt": 2823.52 },
    { "id":19, item: "ipad2", "amt": 623.52 }
  ]
}
```



# Models for Representing Data

Model	Relational Model	JSON Document Model (NoSQL)
1:1	<ul style="list-style-type: none"><li>▪ Foreign Key</li><li>▪ Denormalize</li></ul>	<ul style="list-style-type: none"><li>▪ Embedded Object (implicit)</li><li>▪ Document Key Reference</li></ul>
1:N	<ul style="list-style-type: none"><li>▪ Foreign Key</li></ul>	<ul style="list-style-type: none"><li>▪ Embedded Array of Objects</li><li>▪ Document key Reference</li></ul>
N:M	<ul style="list-style-type: none"><li>▪ Foreign Key</li></ul>	<ul style="list-style-type: none"><li>▪ Embedded Array of Objects, arrays of arrays with references</li></ul>



3

# JSON Design Choices



# JSON Design Choices

---

- Single Root Attributes vs. "type"/"class" parameter
- Objects vs. Arrays
- Array Element Types
- Timestamp Formats
- Property Names
- Empty and Null Property Values
- JSON Schema



# Single Root Attributes

```
{  
  track: {  
    artist: "Paul Lekakis",  
    created: "2015-08-18T19:57:07",  
    genre: "Hi-NRG",  
    id: "3305311F4A0FAAFEABD001D324906748B18FB24A",  
    mp3: https://goo.gl/KgKoR7,  
    ratings: [  
      {  
        created: "2015-08-20T12:24:44",  
        rating: 4,  
        username: "sublimatingraga37014"  
      },  
      {  
        created: "2015-08-21T09:23:57",  
        rating: 4,  
        username: "untillableshowings34122"  
      },  
      {  
        created: "2015-08-21T13:53:34",  
        rating: 3,  
        username: "megacephalousfusty75226"  
      }  
    ],  
    title: "My House",  
    updated: "2015-08-18T19:57:07"  
  }  
}
```

```
{  
  artist: "Paul Lekakis",  
  created: "2015-08-18T19:57:07",  
  genre: "Hi-NRG",  
  id: "3305311F4A0FAAFEABD001D324906748B18FB24A",  
  mp3: https://goo.gl/KgKoR7,  
  ratings: [  
    {  
      created: "2015-08-20T12:24:44",  
      rating: 4,  
      username: "sublimatingraga37014"  
    },  
    {  
      created: "2015-08-21T09:23:57",  
      rating: 4,  
      username: "untillableshowings34122"  
    },  
    {  
      created: "2015-08-21T13:53:34",  
      rating: 3,  
      username: "megacephalousfusty75226"  
    }  
  ],  
  title: "My House",  
  type: "track",  
  updated: "2015-08-18T19:57:07"  
}
```



# Objects vs. Arrays

- Two different ways to represent attributes:

```
{  
  ▶ userprofile: {  
    ▶ address: {...},  
    created: "2015-01-28T13:50:56",  
    dateOfBirth: "1986-06-09",  
    email: "andy.bowman@games.com",  
    ▶ favoriteGenres: [...],  
    firstName: "Andy",  
    gender: "male",  
    lastName: "Bowman",  
    ▶ phones: {  
      cell: "212-771-1834"  
    },  
    ▶ picture: {...},  
    pwd: "636f6c6f7261646f",  
    status: "active",  
    title: "Mr",  
    updated: "2015-08-25T10:29:16",  
    username: "copilotmarks61569"  
  }  
}
```

```
{  
  ▶ userprofile: {  
    ▶ address: {...},  
    created: "2015-01-28T13:50:56",  
    dateOfBirth: "1986-06-09",  
    email: "andy.bowman@games.com",  
    ▶ favoriteGenres: [...],  
    firstName: "Andy",  
    gender: "male",  
    lastName: "Bowman",  
    ▶ phones: [  
      ▶ {  
        number: "212-771-1834",  
        type: "cell"  
      }  
    ],  
    ▶ picture: {...},  
    pwd: "636f6c6f7261646f",  
    status: "active",  
    title: "Mr",  
    updated: "2015-08-25T10:29:16",  
    username: "copilotmarks61569"  
  }  
}
```



# Array Element Types

- Array elements can be simple types, objects or arrays:

```
{  
  "playlist": {  
    "created": "2014-12-04T03:36:18",  
    "id": "003c6f65-641a-4c9a-8e5e-41c947086cae",  
    "name": "Eclectic Summer Mix",  
    "owner": "copilotmarks61569",  
    "tracks": [  
      "9FFAF88C1C3550245A19CE3BD91D3DC0BE616778",  
      "3305311F4A0FAAFEABD001D324906748B18FB24A",  
      "0EB4939F29669774A19B276E60F0E7B47E7EAF58"  
    ],  
    "updated": "2015-09-11T10:39:40",  
    "visibility": "PUBLIC"  
  }  
}
```

Array of strings

```
{  
  "playlist": {  
    "created": "2014-12-04T03:36:18",  
    "id": "003c6f65-641a-4c9a-8e5e-41c947086cae",  
    "name": "Eclectic Summer Mix",  
    "owner": "copilotmarks61569",  
    "tracks": [  
      {  
        "id": "9FFAF88C1C3550245A19CE3BD91D3DC0BE616778"  
      },  
      {  
        "id": "3305311F4A0FAAFEABD001D324906748B18FB24A"  
      },  
      {  
        "id": "0EB4939F29669774A19B276E60F0E7B47E7EAF58"  
      },  
    ],  
    "updated": "2015-09-11T10:39:40",  
    "visibility": "PUBLIC"  
  }  
}
```

Array of objects



# Timestamp Formats

- When storing timestamps, you have at least 3 options:

```
{  
  <country>:  
    <countryCode>: "US",  
    <gdp>: 53548,  
    <name>: "United States of America",  
    <population>: 325296592,  
    <region>: "Americas",  
    <region-number>: 21,  
    <sub-region>: "Northern America",  
    <updated>: "2010-07-15T15:34:27"  
}
```

String (ISO 8601)

```
{  
  <country>:  
    <countryCode>: "US",  
    <gdp>: 53548,  
    <name>: "United States of America",  
    <population>: 325296592,  
    <region>: "Americas",  
    <region-number>: 21,  
    <sub-region>: "Northern America",  
    <updated>: 1279208067000  
}
```

Number (Unix style)

```
{  
  <country>:  
    <countryCode>: "US",  
    <gdp>: 53548,  
    <name>: "United States of America",  
    <population>: 325296592,  
    <region>: "Americas",  
    <region-number>: 21,  
    <sub-region>: "Northern America",  
    <updated>: [  
      2010,  
      7,  
      15,  
      15,  
      34,  
      27  
    ]  
}
```

Array of time components

- Make timestamp values relative to UTC



# Property Names

---

- Choose meaningful property names
- Be consistent in naming properties
  - e.g. country\_code vs. countryCode (preferred)
- Array types should have plural property names
- All other property names should be singular
- Avoid (if possible) reserved words in your database system and programming language(s)
  - e.g. `user` // Reserved word in Couchbase Server
- Avoid (if possible) special characters such as hypens
  - e.g. `region-number` // Contains a hyphen



# Empty and Null Property Values

- Keep in mind that JSON supports optional properties
- If a property has a null value, consider dropping it from the JSON, unless there's a good reason not to
- N1QL makes it easy to test for missing or null property values

```
SELECT * FROM couchmusic1 WHERE userprofile.address IS NULL;
```

```
SELECT * FROM couchmusic1 WHERE userprofile.gender IS MISSING;
```

- Be sure your application code handles the case where a property value is missing



# JSON Schema

- Couchbase Server pays absolutely no attention to the shape of your JSON documents so long as they are well-formed
- There are times when it is useful to validate that a JSON document conforms to some expected shape
- JSON Schema is a JSON-based format for defining the structure of JSON data
- There are implementations for most popular programming languages

```
{
  id: "http://couchmusic.org/schema/couchmusic2-country.json",
  $schema: "http://json-schema.org/draft-04/schema#",
  type: "object",
  properties: {
    countryCode: {
      type: "string",
      minLength: 2,
      maxLength: 2
    },
    gdp: {
      type: "integer",
      minimum: 0
    },
    name: {
      type: "string"
    },
    population: {
      type: "number",
      minimum: 0
    },
    region-number: {
      type: "integer",
      minimum: 0
    },
    type: {
      enum: [
        "country"
      ]
    },
    updated: {
      type: "string",
      format: "date-time"
    }
  },
  required: [
    "countryCode",
    "gdp",
    "name",
    "population",
    "region-number",
    "updated"
  ],
  additionalProperties: false
}
```



# 4

## Data Nesting



# Data Nesting (aka Denormalization)

---

- As you know, relational database design promotes separating data using normalization, which doesn't scale
- For NoSQL systems, we often avoid normalization so that we can scale
- Nesting allows related objects to be organized into a hierarchical tree structure where you can have multiple levels of grouping
- Rule of thumb is to nest no more than 3 levels deep unless there is a very good reason to do so
- You will often want to include a timestamp in the nested data



## Example #1 of Data Nesting

Playlist with owner attribute containing username of corresponding userprofile

```
{  
  > playlist: {  
    created: "2014-12-04T03:36:18",  
    id: "003c6f65-641a-4c9a-8e5e-41c947086cae",  
    name: "Eclectic Summer Mix",  
    owner: "copilotmarks61569",  
    > tracks: [...],  
    updated: "2015-09-11T10:39:40",  
    visibility: "PUBLIC"  
  }  
}
```

```
{  
  > userprofile: {  
    address: {...},  
    created: "2015-01-28T13:50:56",  
    dateOfBirth: "1986-06-09",  
    email: "andy.bowman@games.com",  
    > favoriteGenres: [...],  
    firstName: "Andy",  
    gender: "male",  
    lastName: "Bowman",  
    > phones: {...},  
    > picture: {...},  
    pwd: "636f6c6f7261646f",  
    status: "active",  
    title: "Mr",  
    updated: "2015-08-25T10:29:16",  
    username: "copilotmarks61569"  
  }  
}
```



# Example #1 of Data Nesting

Playlist with owner attribute containing a subset of the corresponding userprofile

```
{  
  > playlist: {  
    created: "2014-12-04T03:36:18",  
    id: "003c6f65-641a-4c9a-8e5e-41c947086cae",  
    name: "Eclectic Summer Mix",  
    > owner: {  
      created: "2015-01-28T13:50:56",  
      firstName: "Andy",  
      lastName: "Bowman",  
      title: "Mr",  
      updated: "2015-08-25T10:29:16",  
      username: "copilotmarks61569"  
    },  
    > tracks: [...],  
    updated: "2015-09-11T10:39:40",  
    visibility: "PUBLIC"  
  }  
}
```

```
{  
  > userprofile: {  
    address: {...},  
    created: "2015-01-28T13:50:56",  
    dateOfBirth: "1986-06-09",  
    email: "andy.bowman@games.com",  
    > favoriteGenres: [...],  
    firstName: "Andy",  
    gender: "male",  
    lastName: "Bowman",  
    > phones: {...},  
    > picture: {...},  
    pwd: "636f6c6f7261646f",  
    status: "active",  
    title: "Mr",  
    updated: "2015-08-25T10:29:16",  
    username: "copilotmarks61569"  
  }  
}
```

\* Note the inclusion of the **updated** attribute



## Example #2 of Data Nesting

# Playlist with tracks attribute containing an array of track IDs

```
{  
  <div><ul>  
    <li>`> playlist: {  
      created: 1417685778000,  
      id: "003c6f65-641a-4c9a-8e5e-41c947086cae",  
      name: "Eclectic Summer Mix",  
      > owner: {...},  
      <div><ul>  
        <li>`> tracks: [  
          "9FFAF88C1C3550245A19CE3BD91D3DC0BE616778",  
          "3305311F4A0FAAFEABD001D324906748B18FB24A",  
          "0EB4939F29669774A19B276E60F0E7B47E7EAF58"  
        ],  
        updated: 1441985980000,  
        visibility: "PUBLIC"  
      </ul></div>  
    </li>  
  </ul></div>
```



## Example #2 of Data Nesting

# Playlist with tracks attribute containing an array of track IDs

```
{  
  ▶ playlist: {  
    created: 1417685778000,  
    id: "003c6f65-641a-4c9a-8e5e-41c947086cae",  
    name: "Eclectic Summer Mix",  
    ▶ owner: {...},  
    ▶ tracks: [  
      ▶ {  
        artist: "Gene Harris",  
        genre: "Jazz Blues",  
        id: "9FFAF88C1C3550245A19CE3BD91D3DC0BE616778",  
        mp3: https://goo.gl/DEYx4X,  
        title: "Battle Hymn of the Republic",  
        updated: 1445167377000  
      },  
      ▶ {...},  
      ▶ {...}  
    ],  
    updated: 1441985980000,  
    visibility: "PUBLIC"  
  }  
}
```

\* Note the inclusion of the **updated** attribute



5

# Key Design



## Natural Keys

---

- A key formed of attributes that exist in the real world:
  - Phone numbers
  - Usernames
  - Social security numbers
  - Account numbers
  - SKU, UPC or QR codes
  - Device IDs
- Often the first choice for document keys
- Be careful when working with any personally identifiable information (PII), sensitive personal information (SPI) or protected health information (PHI)





## Surrogate Keys

---

- We often use surrogate keys when no obvious natural key exist
- They are not derived from application data
- They can be generated values
  - 3305311F4A0FAAFEABD001D324906748B18FB24A (SHA-1)
  - 003C6F65-641A-4CGA-8E5E-41C947086CAE (UUID)
  - They can be sequential numbers (often implemented using the Counter feature of Couchbase Server)
  - 456789, 456790, 456791, ...



# Key Value Patterns

---

- Common practice for users of Couchbase Server to follow patterns for formatting key values by using symbols such as single or double colons
- DocType::ID
  - userprofile::fredsmith79
  - playlist::003c6f65-641a-4c9a-8e5e-41c947086cae
- AppName::DocType::ID
  - couchmusic::userprofile::fredsmith79
- DocType::ParentID::ChildID
  - playlist::fredsmith79::003c6f65-641a-4c9a-8e5e-41c947086cae
- Supports easy document viewing in the Couchbase web console



# Lookup Key Pattern

---

- The purpose of the Lookup Key Pattern is to allow multiple ways to reach the same data, essentially a secondary index
  - For example, we want to lookup a Userprofile by their email address instead of their ID
- To accomplish this, we create another small document that refers to the Userprofile document we are interested in
- Implementing this pattern is straightforward, just create an additional document containing a single property that stores the key to the primary document
- With the introduction of N1QL, this pattern will be less commonly used

# Example of Lookup Key Pattern

- Lookup document can be JsonDocument or StringDocument

userprofile::copilotmarks61569

```
{  
  <--> userprofile: {  
    > address: {...},  
    created: "2015-01-28T13:50:56",  
    dateOfBirth: "1986-06-09",  
    email: "andy.bowman@games.com",  
    > favoriteGenres: [...],  
    firstName: "Andy",  
    gender: "male",  
    lastName: "Bowman",  
    > phones: {...},  
    > picture: {...},  
    pwd: "636f6c6f7261646f",  
    status: "active",  
    title: "Mr",  
    updated: "2015-08-25T10:29:16",  
    username: "copilotmarks61569"  
  }  
}
```

andy.bowman@games.com

```
{  
  <--> username: "copilotmarks61569"  
}
```

JSON

andy.bowman@games.com

```
copilotmarks61569
```

String



6

# Making Trade-Offs



# Making Tough Choices

---

- Eric Brewer is famous for showing the trade-offs that are necessary when dealing with distributed systems
  - Consistency, availability and partition tolerance are all desirable properties but we must choose the ones that are most important for our use cases
- We must also make trade-offs in data modeling:
  - Document size
  - Complexity
  - Speed
  - Atomicity





## Document Size

---

- Couchbase Server supports documents up to 20 Mb
- Larger documents take more disk space, more time to transfer across the network and more time to serialize/deserialize
- If you are dealing with documents that are potentially large (greater than 1 Mb), you must test thoroughly to find out if speed of access is adequate as you scale. If not, you will need to break up the document into smaller ones.
- You may need to limit the number of dependent child objects you embed



# Complexity

- Complexity affects every area of software systems including data modeling
- Space saving sometimes comes at the complexity of data processing
- E.g., consider the complexity of queries (N1QL)

```
SELECT c.name, COUNT(*) AS playlist_count
FROM couchmusic1.playlist p JOIN couchmusic1.userprofile u
    ON KEYS 'userprofile:::' || p.owner JOIN couchmusic1.country c
        ON KEYS 'country:::' || u.address.countryCode
WHERE c.`region-number` = 154
    AND p.visibility = 'PUBLIC'
    AND u.status = 'active'
GROUP BY c.name
ORDER BY c.name;
```



# Speed

---

- As it relates to data modeling, speed of access is critical
- When using N1QL to access data, keep in mind that query by document key is fastest and query by secondary index is usually much slower
- If implementing an interactive use case, you will want to avoid using JOINS
- You can use data duplication to improve the speed of accessing related data and thus trade improved speed for greater complexity and larger document size
- Keep in mind that Couchbase Views can be used when up to the second accuracy is not required



# Atomicity

---

- Atomicity in Couchbase Server is at the document level
- Couchbase Server does not support transactions
  - They can be simulated if you are willing to write and maintain additional code to implement them (generally not recommended)
- If you absolutely need changes to be atomic, they will have to be part of the same document
- The maximum document size for Couchbase Server may limit how much data you can store in a single document
- Multi-document transactions will be introduced in Couchbase 6.5



# Embed vs. Refer

---

- When to embed:
  - Reads greatly outnumber writes
  - You're comfortable with the slim risk of inconsistent data across the multiple copies
  - You're optimizing for speed of access
- When to refer:
  - Consistency of the data is a priority
  - You want to ensure your cache is used efficiently
  - The embedded version would be too large or complex





# Summary

---

- In this module, you have learned to:
  - Make full use of JSON capabilities
  - Use data nesting to minimize the need for JOINs
  - Establish key value patterns and use them consistently
  - Be clear about the trade-offs you are making, document your decisions and the assumptions they are based on



# Thank you



Couchbase