



Couchbase

# What's new in Couchbase 6.6



# Enterprise-Class NoSQL Database

	CB 6.0	New in CB 6.5 & 6.6
<b>ACID Transactions</b>	Single-Document ACID	Multiple-Document ACID
<b>Query Enhancements</b>	SQL for JSON (Query & Analytics), Declarative, Visual Query Plan, Search	ANSI Window Functions, ANSI CTE, Query+Search, Cost-Based Optimizer (Preview) Index Advisor (Preview) (GA in 6.6)
<b>Eventing</b>	Server-side programming	CURL, Source Bucket Mutation, User-Defined Functions (Preview)
<b>SDK</b>	SDKs with support for all Couchbase services	Scala Support, Simplified Async Programming, Improved HA & Observability, Improved APIs
<b>High App &amp; Data Density</b>	10 Buckets	30 Buckets, Collections (Preview), High Data Density (Private Preview)
<b>HA/DR</b>	Active-Active, Memory-Memory Streaming	XDCR - Quality of Service, Advanced Filtering, Robust Rebalance, RMS for GSI
<b>Backup/Restore</b>	Full & Incremental Backup/Restore	60% better throughput, reduce storage requirements by 2x-4x
<b>Security &amp; Compliance</b>	X.509 Authentication, RBAC, Secure Transport, Auditing, Field-Level Encryption	LDAP Group, Node-Node Encryption, TLS Cipher Settings

# **1** Distributed ACID Transactions



# Distributed ACID Transactions

Simplify application logic with distributed ACID transactions, providing all-or-nothing semantics for durably modifying multiple documents distributed on different nodes and across buckets.

Transition from RDBMS schema:

- De-normalization has limitations

Multi-Asset Coordination:

- Transfer from one user to another
- Reservation items such as flights

Business/Application-level transaction needs to modify multiple documents (all-or-nothing):

- Microservices SAGAs - Event based orchestration

```
transactions.run((txnctx) -> {
    // Insert a document
    JsonDocument doc1 =
        JsonDocument.create("newDoc", JsonObject.create());
    txnctx.insert(bucket, doc1);

    // Replace a document
    TransactionJsonDocument doc2 =
        txnctx.getOrError(bucket, "doc2");
    doc2.content().put("name", "bob");
    txnctx.replace(doc2);

    // Commit transaction
    txnctx.commit();
});
```



# ACID in Couchbase

Couchbase Server 6.5 provides strong ACID guarantees while balancing scalability, availability and performance.

A	Atomicity	Guarantees all-or-nothing semantics for <b>updating multiple documents</b> in more than one shards on different nodes.
C	Consistency	<b>Replicas</b> immediately consistent with the transaction commit. <b>Indexes and XDCR</b> eventually consistent with the transaction commit ( <i>N1QL can enforce strong consistency upon read with request_plus</i> )
I	Isolation	<b>Read Committed</b> isolation for concurrent transactions
D	Durability	<b>Data protection under failures</b> : 3 different levels - <i>replicate to majority of the nodes; replicate to majority and persist to disk on primary; or persist to disk on majority of the nodes.</i>



# Couchbase Transactions: Built on Synchronous Replication

Synchronous Replication ensures that a mutation is visible to readers only after it has met its replication/persistence criteria and can be rolled back if those criteria are not met.

- SyncWrite is **all-or-nothing**
- **Sequential Consistency model**
- **3 Durability levels** (per-write)
  - Replicate to Majority
  - Replicate to Majority and Persist to Master
  - Persist to Majority
- Applicable to **transactions and single document mutations**
- Couchbase and Ephemeral buckets
- In-house **Jepsen** testing

```
MutationResult result =
    collection.insert(id,
new Person("Jon Henry", 34, null, null),
insertOptions()
    .durabilityLevel(DurabilityLevel.
        MAJORITY_AND_PERSIST_ON_MASTER)
    .timeout(Duration.ofSeconds(30)))
    );
```



# Couchbase Transactions: ACID at Scale with High Availability

- **No SPOF or performance bottleneck**
  - Couchbase Smart Clients act as the Transaction Coordinators
    - No reliance on a global transaction coordinator
  - Chosen algorithms ensure progress only depends upon availability of shards involved in the transactions
    - Read Committed: no global scheduler
    - Built-in retries: no distributed lock manager
- **Pay the transaction overhead only when you use it**
  - Flexibility to use a transaction is available at the level of an operation
  - No need to unconditionally impose it for every operation

# **2** Query Enhancements



# Query Enhancements

**Window  
Functions**

**Common  
Table  
Expressions**

**SEARCH in  
N1QL**

**User Defined  
Functions**  
(preview)

**Index Advisor**  
(GA in 6.6)

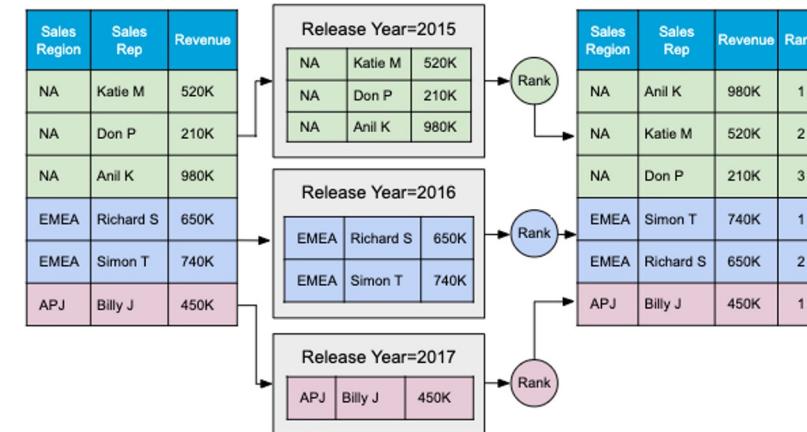
**Cost-Based  
Optimizer**  
(GA in 7.0)



# ANSI Window Functions

With ANSI Window Functions, developers can simplify financial and statistical aggregations in an easy and optimized way

- Answer common (but complex) business queries with **minimal lines of code** and **optimized performance**
- Use cases:
  - revenue growth month over month
  - top N sale districts by revenues for a given week
  - ranking of sales person by region based on revenue booked
- Couchbase is the first NoSQL Database to support ANSI Window Functions



```
SELECT sales_region, release_year, sales_rep,  
       RANK() OVER (PARTITION BY sales_region, release_year  
                      ORDER BY SUM(revenue) DESC) as `rank`  
FROM orders  
GROUP BY sales_region, sales_rep, release_year
```



# ANSI Common Table Expression (CTE)

With ANSI Common Table Expressions, developers have ease of maintenance and better readability by naming temporary SQL statements.

- CTE allows developer to isolate SQL statement into **temporary named result set** that can be referenced as a source table in the context of a larger query
- Offers the advantages of **readability and ease of maintenance** of complex queries without compromising performance.
- Couchbase is the first NoSQL Database to support ANSI CTE.

```
WITH current_period_task AS (
    SELECT DATE_TRUNC_STR(a.startDate,'month') AS month,
    COUNT(1) AS current_period_task_count
    FROM crm a
    WHERE a.type='activity' AND a.activityType = 'Task'
    AND DATE_PART_STR(a.startDate,'year') = 2018
    GROUP BY DATE_TRUNC_STR(a.startDate,'month')
),
last_period_task AS (
    SELECT x.month, x.current_period_task_count,
    LAG(x.current_period_task_count) OVER (ORDER BY x.month)
    AS last_period_task_count
    FROM current_period_task x
)
SELECT b.month,
    b.current_period_task_count,
    ROUND(((b.current_period_task_count - b.last_period_task_count) /
    b.last_period_task_count),2)
FROM last_period_task AS b
```

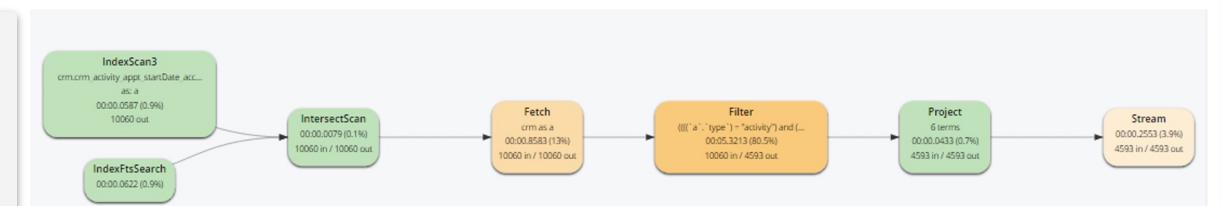


# SEARCH Predicates in Query (Full Text Search)

With Search predicates in N1QL, developers can combine SQL and Search queries for powerful integration and better query performances

- Couchbase provides a comprehensive search capability in query **beyond the simple LIKE() operator** in most databases
- The **SEARCH()** operator supports keyword and fuzzy matchings across multiple document fields
- Developers **do not have to write complex code** to process and combine the results from separate SQL and search queries
- **Better query performance** with inverted indexes for search predicates instead of inefficient scans for LIKE()

```
SELECT * FROM `beer-sample` b
WHERE SEARCH(b.desc, "fruity")
AND b.abv < 0.5;
```





# Index Advisor

With Index Advisor, developer can create better indexes based on the suggestions to speed up queries easily.

- The Index Advisor **suggests appropriate indexes** to speed up a given query, taking the guesswork out of query tuning
- The Index Advisor can also **monitor and analyze the statistics** collected from running a workload, and suggest the indexes that will speed up the queries in the workload
- The Index Advisor greatly **reduces the complexity and efforts** required for enterprise developers and operations engineers to determine the right indexes to speed up their queries

Query Editor < history (551/551) >

```
1 SELECT *
2 FROM `beer-sample`
3 WHERE `type` = 'beer'
```

Execute Explain Advise executing

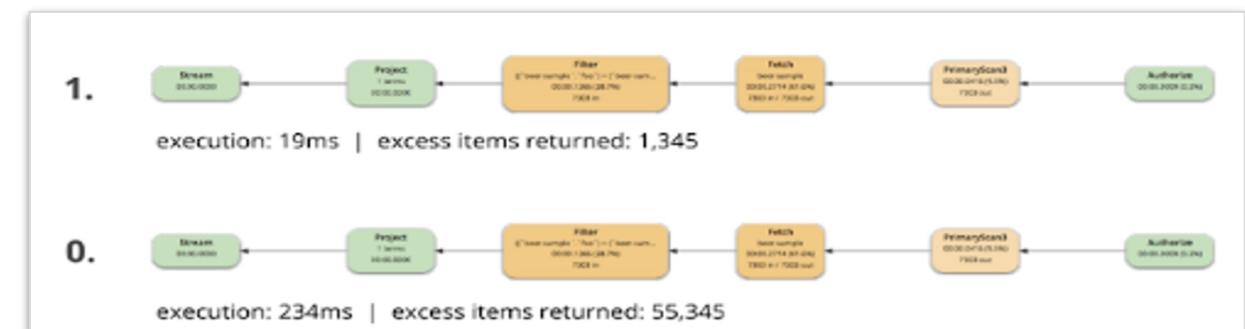
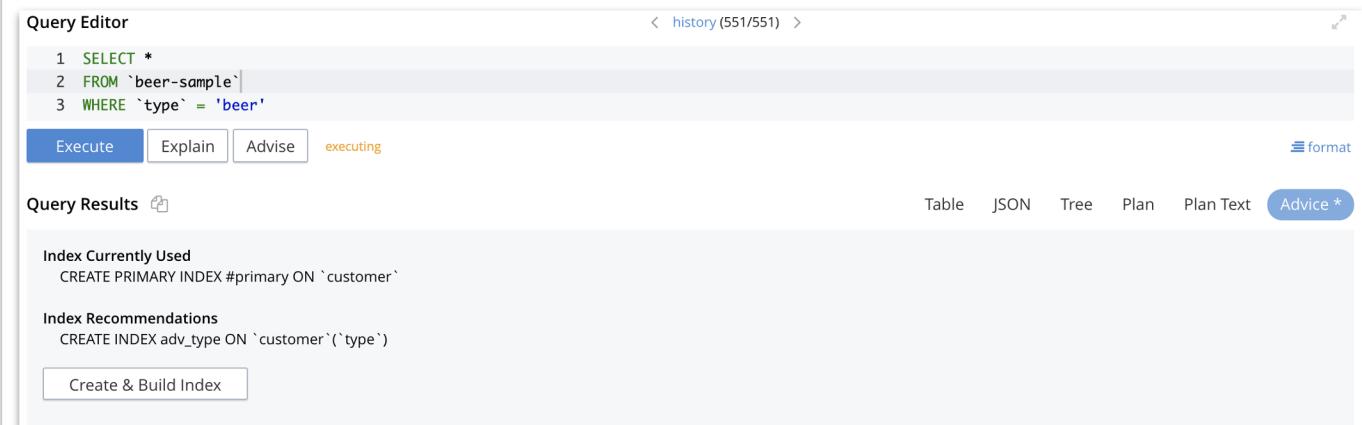
Query Results

Index Currently Used  
CREATE PRIMARY INDEX #primary ON `customer`

Index Recommendations  
CREATE INDEX adv\_type ON `customer`(`type`)

Create & Build Index

Table JSON Tree Plan Plan Text Advice \*





# User-Defined Functions (Developer Preview)

With User-Defined Functions, developers define callable custom functions to ease development, simplify reusability and improve code performance.

- Allow developers to define **custom functions** in N1QL or Javascript (similar to PL/SQL) callable from N1QL queries.  
Interactive debugger to simplify development and testing
- Server-side logic that can be **reused** by any application and micro services; improve code maintenance and developer productivity
- Improve code performance** by bringing application logic closer to the data

```
CREATE FUNCTION to_meters() { args[0] / 0.3048 };
```

```
SELECT name, ROUND(to_meters(geo.alt)) AS mams1  
FROM `travel-sample`  
WHERE type = "airport"  
LIMIT 1;
```

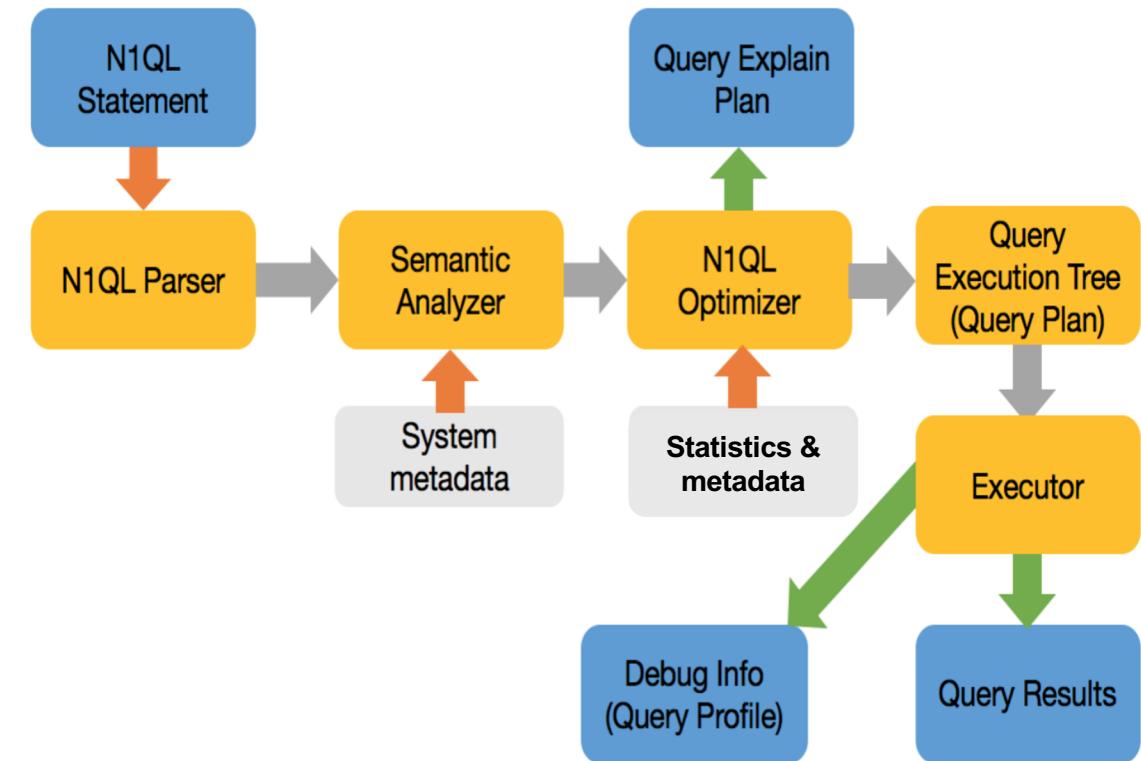
```
{  
  "airportname": "Calais Dunkerque",  
  "mams1": 39  
},
```



# Cost-Based Optimizer (Developer Preview)

With Cost-Based Optimizer, queries will run faster by using optimal access path without the need to provide optimizer hints.

- Cost-based optimizer that generates the **optimal access path based on statistics** collected on the data
- Eliminates time tweaking a query and providing optimizer hints to get the rule-based optimizer to pick the right query plan
- Leverages decades of research and experience in query optimization to collect and use statistics on JSON, arrays, and objects.
- Couchbase is the first NoSQL database with dynamic schema to support cost-based optimizer





# Easy Migration from RDBMS to Couchbase

## MySQL

```
SELECT ac.industry,
       SUM(CASE WHEN a.activitytype = 'Task'
              THEN 1 ELSE 0 END ) task,
       SUM(CASE WHEN a.activitytype = 'Appointment'
              THEN 1 ELSE 0 END ) appts
  FROM crm.activity a
 INNER JOIN crm.account ac
    ON (a.accid = ac.id)
 WHERE a.startdate BETWEEN '2018-10-01'
      AND '2018-12-31'
 GROUP BY ac.industry
```

## Couchbase - N1QL

```
SELECT ac.industry,
       SUM(CASE WHEN a.activityType = 'Task'
              THEN 1 ELSE 0 END) task,
       SUM(CASE WHEN a.activityType = 'Appointment'
              THEN 1 ELSE 0 END) appts
  FROM crm a
 INNER JOIN crm ac ON a.accid = ac.id
          AND ac.type='account'
 WHERE a.type='activity'
       AND a.startDate BETWEEN '2018-10-01'
      AND '2018-12-31'
 GROUP BY ac.industry
```

## MongoDB Query Language

```
db.activity.aggregate(
  { $match: { startDate: { $gt: '2018-01-01',
                         $lt: '2018-12-31' } } },
  {
    $lookup: {
      from: "account",
      localField: "accid",
      foreignField: "id",
      as: "account_docs"
    }
  },
  { $match: { "account_docs": { $ne: [] } } },
  { $unwind: "$account_docs" },
  {
    $project: {
      item: 1,
      task: { $cond: { if:
                      { $eq: ["$activityType", "Task"] }, then: 1,
                      else: 0 } },
      appt: { $cond: { if:
                      { $eq: ["$activityType", "Appointment"] }, then: 1,
                      else: 0 } } }
  },
  {
    $group: {
      _id: "$account_docs.industry",
      tasks: { $sum: "$task" },
      appointments: { $sum: "$appt" }
    }
  }
);
```

# 3 SDK 3.0



**Scala**

**Simplified  
Async  
Programming**

**Improved  
HA and  
Observability**

**Improved  
APIs**



# New Platform: Scala



- Native Scala APIs - get and upsert documents, run queries, perform analytics and full text searches
- Has built-in support for JSON and also integrates with popular Scala JSON libraries: Circe, μPickle / μJson, Json4s and Jawn
- Supports both Scala Futures and Project Reactor based reactive programming with **Mono and Flux**

```
...
cluster = Cluster.connect(env)
val bucket = cluster.bucket(config.bucketname)
...
val content = ujson.Obj("hello" -> "world")
val upsertResult = coll.upsert(docId, content)

upsertResult match {
  case Success(result) =>
    // do something interesting
  case Failure(err) =>
    // handle some errors
}
val id: String = insertDoc
for (i <- Range(0, 10)) {
  val options: QueryOptions = QueryOptions()
    .scanConsistency(ScanConsistency.RequestPlus())
    .adhoc(false)
    .namedParameters("id" -> id)
  val result: QueryResult = cluster.query("select `"
+ bucketName + `.* from ` + bucketName + ` where
meta() +
  ".id=$id", options).get
  val rows = result.allRowsAs[JsonObject].get
```



# Simplified Async Programming

Simplifies asynchronous programming and avoids problems associated with callbacks, leveraging the latest async frameworks/concepts in key languages

- Node.js SDK - Supports Promises in JavaScript ES6
- Java SDK - Supports Reactor on top of reactive streams and supports Java 8 CompletableFuture
- Scala SDK - Supports Reactor, Scala Native Futures, Akka Streams
  - Opens Couchbase up to the Akka developer ecosystem out of the box
- Python SDK - Supports Async IO in Python 3



# Simplified Async Programming

```
function attemptChange(callback) {
  bucket.get('somedoc', function(err, res) {
    if (err) {
      callback(err);
      return;
    }

    if (res.value > 100) {
      res.value += 1;
    } else {
      res.value += 10;
    }

    bucket.replace('somedoc', res.value, {cas: res.cas}, function(err, res) {
      if (err && err.code == 0x02) {
        // CAS mismatch on upsert, try again
        attemptChange(callback);
        return;
      } else if (err) {
        callback(err);
        return;
      }
    })
  })
}

attemptChange(function(err) {
  if (err) {
    console.log('Change Errorred', err);
  } else {
    console.log('Change Completed');
  });
})
```



# Simplified Async Programming

- No more nested callbacks!
- Performs like asynchronous code, but looks simpler
- Simplified async error handling with await keyword from ECMAScript 6.x+

```
do {
    var res = await coll.get('somedoc');

    if (res.value > 100) {
        res.value += 1;
    } else {
        res.value += 10;
    }
    try {
        await coll.replace('somedoc', res.value, {cas: res.cas});
    } catch (e) {
        if (e instanceof couchbase.CasMismatchError) {
            continue;
        }

        // rethrow real errors.
        throw e;
    }
} while(false);
```



# Improved HA and Observability

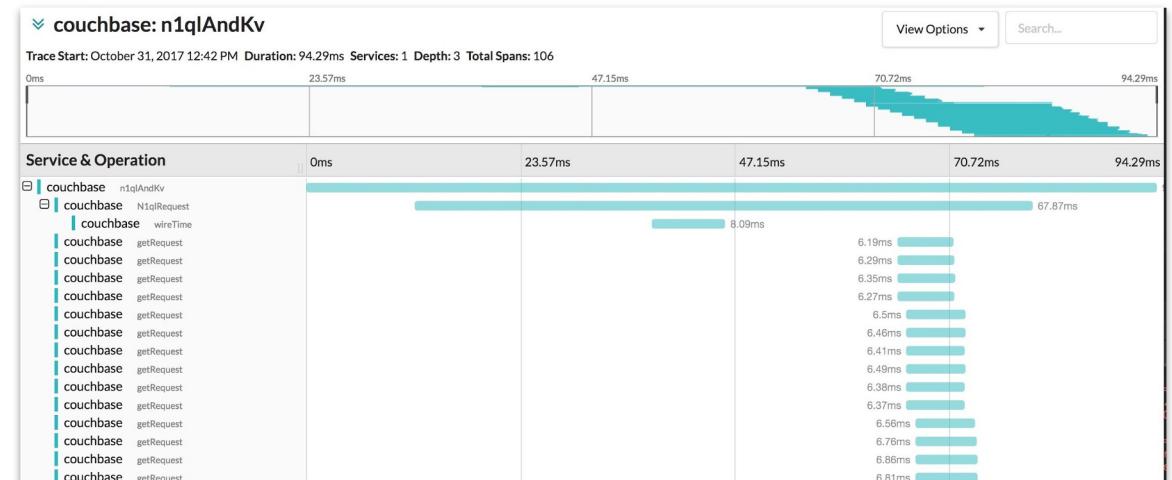
## Circuit Breaker

- Improve system availability and prevent cascading failures by avoiding putting loads on a struggling resource
- Provides early warning on potential system issues when circuit breaker trips.

```
CircuitBreakerConfig.builder()  
    .enabled(true)  
    .errorThresholdPercentage(50)  
    .rollingWindow(Duration.ofSeconds(60))  
    .sleepWindow(Duration.ofSeconds(5))  
    .volumeThreshold(20)  
    .build();
```

## OpenTelemetry

- Make it easy to get critical telemetry data from Couchbase applications in a robust and portable way
- OpenTelemetry is a CNCF sandbox project, and the next major version of the OpenTracing and OpenCensus projects





# Improved APIs

## Reduce Overloads

```
bucket.upsert  
upsert(D document)  
upsert(D document, PersistTo persistTo)  
upsert(D document, ReplicateTo replicateTo)  
upsert(D document, long timeout, TimeUnit timeUnit)  
upsert(D document, PersistTo persistTo, ReplicateTo, rep...  
upsert(D document, PersistTo persistTo, long timeout, ...  
upsert(D document, ReplicateTo replicateTo, long time...)  
...
```

```
MutationResult mutationResult = collection.upsert(Id, person,  
    upsertOptions().expiry(Duration.ofDays(1))  
);
```

## Prepare for Collections

- Old: Cluster -> Bucket -> Document
- New: Cluster -> Bucket -> Scope -> Collections

```
1 //Open a Scope and then a Collection  
2 var scope = await bucket.Scope("ACME");  
3 var collection = await scope.Collection("tools");  
4  
5 //add a new documentawait  
6 collection.Insert(key, new Tool{Name = "Hammer", Type="Ball peen"});  
7  
8 //retrieve the document  
9 using (var result = await collection.Get(key)) {  
10     var content = result.ContentAs<Tool>();  
11     Console.WriteLine(content.Name);  
12 }
```

# 4 High Density



# High Application and Data Density Enhancements

**30 Buckets**

**Collections**  
(preview)

**HiDD**  
(private preview)



# High Application Density - More Buckets per Cluster

User can create up to 30 buckets in a single Couchbase cluster and operate it seamlessly as long as sizing is done appropriately.

- **Multi-tenancy** support
  - Buckets provide physical, logical and access isolation
  - Improve resource utilization of hardware by co-locating multiple tenants
- **Microservices** consolidation
  - Buckets provide physical, logical and access isolation
  - Improve resource utilization of hardware by co-locating multiple microservices
- Previous limit was up to 10 buckets
- **Sizing still critically important**
  - Factor in raw overhead per bucket
  - Eventing service not supported with multiple functions per bucket if using large number of buckets



# High Application Density - Collections (Dev Preview)

With Collections, application developers can organize and access control their data as separate logical containers within a bucket.

- Logical Data Container within a bucket
- Collections can be organized in Scopes
- Document IDs unique within a collection
- 1000s of collections in a bucket
- Security Isolation via RBAC
- Easier Migration from RDBMS
  - Closer mapping to relational schema
  - N1QL easier and more intuitive
- Improve hardware utilization
  - Consolidation of applications/microservices
  - Soft co-operative multi-tenancy

- **Available in 6.5 DP:**
  - Data Service functionality
  - Cluster manager functionality for collections including REST APIs, Rebalance and Failover handling
  - SDK APIs for collection access
- **GA in 7.0:**
  - RBAC at collection and scope level
  - UI awareness of collections
  - XDCR and Backup/restore at collection level
  - Indexing/Query/FTS/Analytics/Eventing at collection level



# High Data Density - 30 TB per Node (GA in 7.1- Q1 2022)

With High Data Density support, enterprises can have dense data nodes with 30 TB per node with less than 1% memory residency for better TCO

- Fewer nodes in cluster
  - Support High Data Density per node - 30 TB data per node
    - Support with low memory residency <1% memory residency
- Faster disk access performance
  - Lower read response time
  - Higher write throughput
- Less SSD wear out
  - Lower write amplification
  - Compaction should not rewrite large parts of data
- Reduction in required disk capacity allocation
  - Lower space amplification

# 5 HA/DR/Backup



# HA/DR Enhancements

XDCR filtering

XDCR QoS

Backup  
Performance

Backup Size



# XDCR - Advanced Filtering

XDCR advanced filtering enables customers to filter and replicate data based on keys, values and extended metadata to simplify partial replication between data centers.

- Documents to be replicated can be **filtered based on keys, values, meta-data** or any combinations of those.
- **Negative filters** such as `(?!regex),(?<!regex)` will be supported which can be applied to filter all documents except a specific subset.
- **Filtering out TTL, expirations and/or deletes** to create hot/cold datasets.
- **Modify filters at runtime.**

Edit Replication X

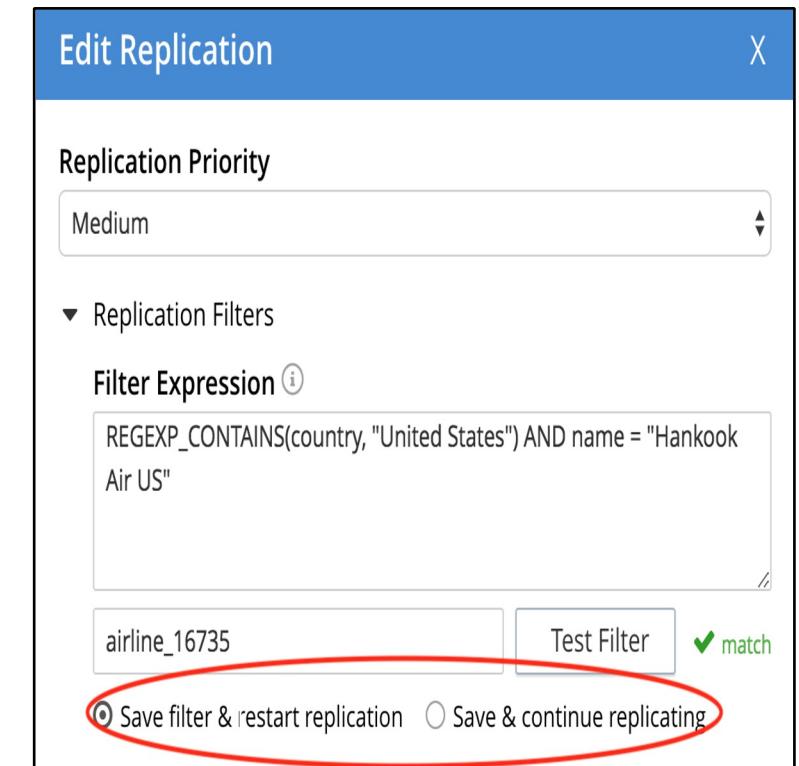
Replication Priority  
Medium

▼ Replication Filters

Filter Expression ⓘ  
`REGEXP_CONTAINS(country, "United States") AND name = "Hankook Air US"`

airline\_16735  ✓ match

Save filter & restart replication  Save & continue replicating





# XDCR - Quality of Service

XDCR quality of service allows control on replication streams by providing an option to prioritize ongoing replications over initial replications to improve uptime.

- Remove limitation of current behavior which always prioritizes new replication over existing ones, sometimes causing degradation of existing streams when new one is created.
- Now possible **to classify a new or ongoing replication as “High, Medium, Low”**, resources allocated accordingly.

Add Replication X

Replicate From Bucket  
travel-sample

Remote Cluster  
test

Remote Bucket  
demo

Low  
Medium  
 High

► Replication Filters  
► Advanced Replication Settings

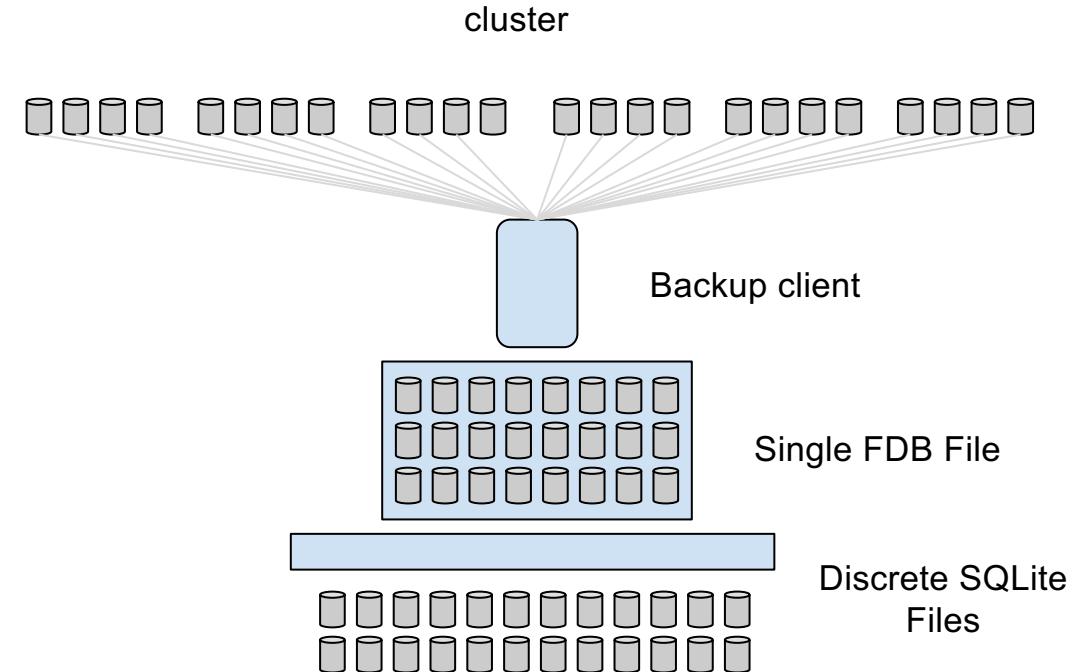
Cancel Save



# Backup - Enhancements

Significantly reduced storage requirement for backup (~40% of original dataset) and performance improvement of more than 60% for ease of administration and better TCO

- Moved from a single monolithic file to 100MB files, introduced DCP compression, replaced ForestDB with SQLite as storage engine and provided the ability to measure consistency.
- Improved List, Merge, Compact performance.
- Auto-bucket creation for restore.
- New Info command (developer preview).



# 6 Manageability Enhancements



# Manageability Enhancements

**Custom  
Statistics**

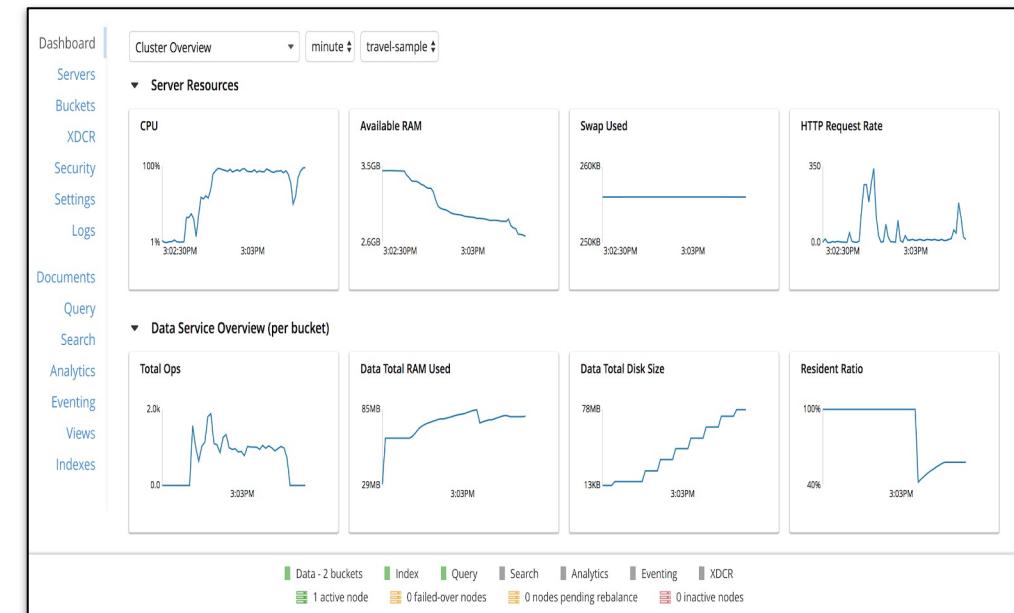
**Robust  
Rebalance**



# Custom UI Statistics

Customizable UI statistics to assist with troubleshooting and simplify monitoring of the whole cluster

- Annotated and labeled stats visuals, responsively sized for any screen
- Ability to use a click-and-drag focus tool to select a custom range of data — zooming in on areas of interest.
- Ability to aggregate stats from multiple nodes to compare and contrast.
- Service specific rendering.
- Ability to edit, group, drop, drag or add the charts
- Users can save the charts created by them.





# Robust Rebalance

Rebalance is now more robust and resilient with automatic restart, better performance, and enhanced monitoring of rebalance progress.

- Auto-restart rebalance on failure
  - Configurable retry interval/count
- Auto-failover during rebalance
  - Prioritize availability
- Improved progress monitoring
  - Status for all services and stages
  - New look and feel
- Lower impact to front-end workloads
  - Scheduling improvements for lower resource utilization

**Rebalance**

start 31 Jul 14:02:07 status in progress

Service	Elapsed	Status
Data	elapsed 00:16	vbuckets moved: 0 of 106 0%
beer-sample		vbuckets moved: 106 of 106 100%
travel-sample		
Analytics	---	
Eventing	---	
Search	---	
Index	---	
Query	---	

**Stop**

# 7 Security Enhancements



# Security Enhancements

**LDAP Groups**

**Node-Node  
Encryption**

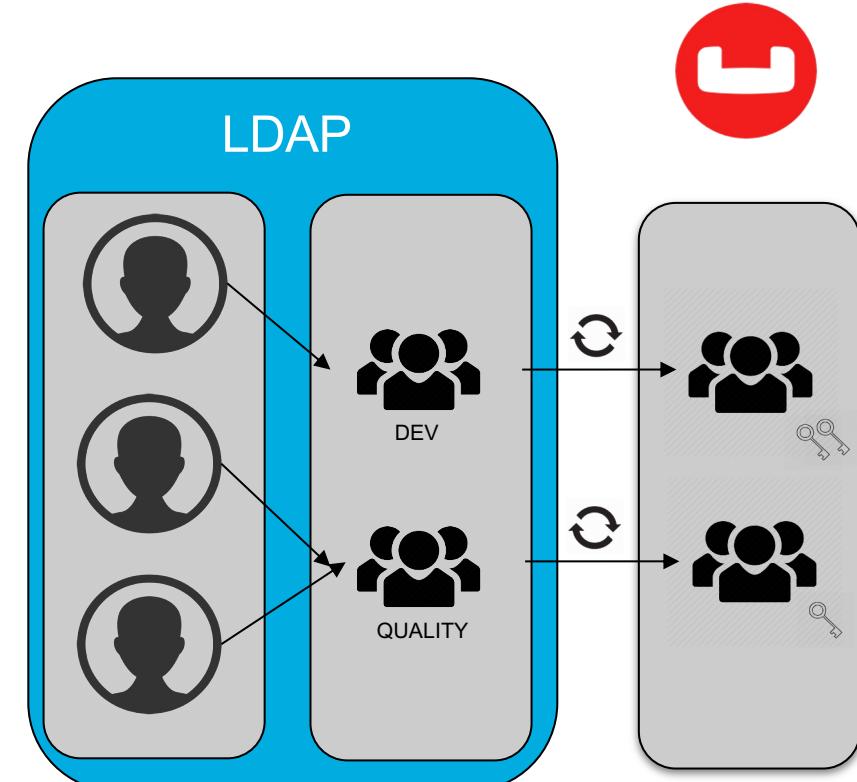
**TLS Cipher  
Management**



# Security - Native LDAP Group Support

LDAP Group support simplifies LDAP integration and includes local group support

- LDAP Group Support enables easy integration with existing LDAP servers, and RBAC user management
- Native LDAP works without SASLauthd support, lighting up LDAP for Couchbase on Windows
- LDAP users and groups reside externally in LDAP
- Couchbase Groups map and sync to LDAP groups
- RBAC privileges assigned to Couchbase groups

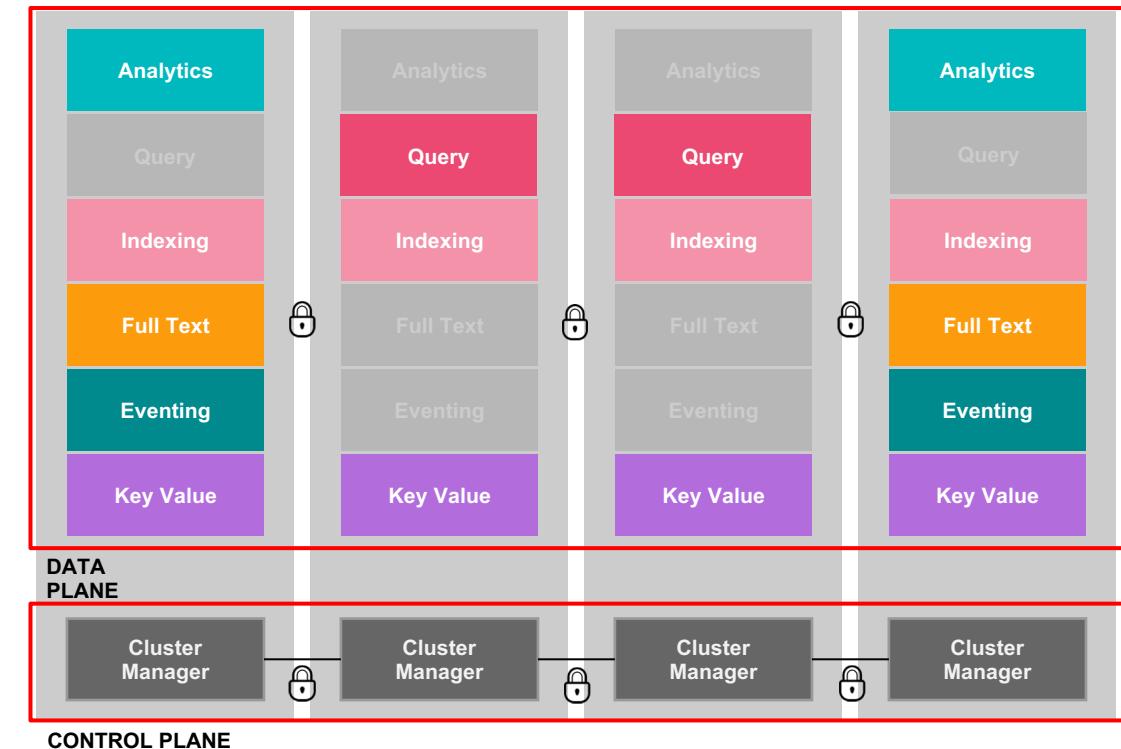




# Security - Node-to-Node Encryption

Node-to-node certificate-based encryption greatly and easily improves security within a cluster.

- On-the-wire encryption for data and control plane, across multiple nodes
- Certificate-based encryption management using existing node and cluster certificates
- Encryption can be controlled and enabled just for the control plane separately
- Using IPSec is not required
- Simple to setup and manage





# Security - TLS Cipher Settings

TLS Cipher settings with centralized management and persistent settings to manage TLS security easily

- **Centralized management of TLS cipher settings:**
  - Get the set of supported ciphers supported across the different Couchbase services
  - Set the list of ciphers to be used consistently across services or per service
  - **Persistent TLS cipher settings** across cluster, node restarts and upgrades

```
curl localhost:8091/settings/security \
-u Administrator:password \
-d 'cipherSuites="[TLS_RSA_WITH_AES_128_CBC_SHA, \
TLS_RSA_WITH_AES_256_CBC_SHA, ...]"'
```



**Couchbase**

# Thank You