



# N1QL



# What is N1QL?

## ***Non-first (N1) Normal Form Query Language (QL)***

*It is based on ANSI 92 SQL*

*Its query engine is optimized for modern, highly parallel multi-core execution*

## **SQL-like Query Language**

Expressive, familiar, and feature-rich language for querying, transforming, and manipulating JSON data

N1QL extends SQL to handle data that is:

**Nested:** Contains nested objects, arrays

**Heterogeneous:** Schema-optional, non-uniform

**Distributed:** Partitioned across a cluster




***Power of  
SQL***

***Flexibility  
of JSON***

# SQL



## *Input: Relations*



## *Output: Relation*

<b>ResultSet</b>	



# N1QL is Declarative: What Vs How

You specify **WHAT**  
Couchbase Server figures  
out **HOW**

## Input: JSON

```
{
  "Name": "Jane Smith",
  "DOB": "1990-01-30",
  "Billing": [
    {
      "type": "visa",
      "cardnum": "5827-2842-2847-3909",
      "expiry": "2019-03"
    },
    {
      "type": "master",
      "cardnum": "6274-2842-2847-3909",
      "expiry": "2019-03"
    }
  ],
  "Connections": [
    {
      "CustId": "XYZ987",
      "Name": "Joe Smith"
    },
    {
      "CustId": "PQR823",
      "Name": "Dylan Smith"
    },
    {
      "CustId": "PQR823",
      "Name": "Dylan Smith"
    }
  ],
  "Purchases": [
    { "id": 12, item: "mac", "amt": 2823.52 },
    { "id": 19, item: "ipad2", "amt": 623.52 }
  ]
}
```

```
{
  "Name": "Jane Smith",
  "DOB": "1990-01-30",
  "Billing": [
    {
      "type": "visa",
      "cardnum": "5827-2842-2847-3909",
      "expiry": "2019-03"
    },
    {
      "type": "master",
      "cardnum": "6274-2842-2847-3909",
      "expiry": "2019-03"
    }
  ],
  "Connections": [
    {
      "CustId": "XYZ987",
      "Name": "Joe Smith"
    },
    {
      "CustId": "PQR823",
      "Name": "Dylan Smith"
    },
    {
      "CustId": "PQR823",
      "Name": "Dylan Smith"
    }
  ],
  "Purchases": [
    { "id": 12, item: "mac", "amt": 2823.52 },
    { "id": 19, item: "ipad2", "amt": 623.52 }
  ]
}
```

{ N1QL }

## Output: JSON

```
{
  "Name": "Jane Smith",
  "DOB": "1990-01-30",
  "Billing": [
    {
      "type": "visa",
      "cardnum": "5827-2842-2847-3909",
      "expiry": "2019-03"
    },
    {
      "type": "master",
      "cardnum": "6274-2842-2847-3909",
      "expiry": "2019-03"
    }
  ],
  "Connections": [
    {
      "CustId": "XYZ987",
      "Name": "Joe Smith"
    },
    {
      "CustId": "PQR823",
      "Name": "Dylan Smith"
    },
    {
      "CustId": "PQR823",
      "Name": "Dylan Smith"
    }
  ],
  "Purchases": [
    { "id": 12, item: "mac", "amt": 2823.52 },
    { "id": 19, item: "ipad2", "amt": 623.52 }
  ]
}
```



# ***N1QL (POWERFUL)***

---

***Access to every part of JSON document***

***JOINS, Aggregations, standard scalar functions***

***Aggregation on arrays***

***NEST & UNNEST operations***

***Covering Index***

***Give developers and enterprises an expressive, powerful, and complete language for querying, transforming, and manipulating JSON data.***

# ***N1QL (QUERYING)***

---



***INSERT***

***UPDATE***

***DELETE***

***MERGE***

***SELECT***

***EXPLAIN***

***Give developers and enterprises an expressive, powerful, and complete language for querying, transforming, and manipulating JSON data.***



# ***N1QL (TRANSFORMING & MANIPULATING)***

***Full Transformation of the data via Query.***

***INSERT***

***INSERT single & multiple documents***

***INSERT result a SELECT statement***

***DELETE documents based on complex filter***

***UPDATE any part of JSON document & use complex filter.***

***MERGE two sets of documents using traditional MERGE statement***

***SUBQUERIES***

***Give developers and enterprises an expressive, powerful, and complete language for querying, transforming, and manipulating JSON data.***



# N1QL (EXPRESSIVE)

---

Access to every part of JSON document

Scalar & Aggregate functions

Subqueries in the FROM clause

Aggregation on arrays

***Expressive, familiar, and feature-rich language for querying, transforming, and manipulating JSON data***





# N1QL (FAMILIAR)

---

SELECT \* FROM bucket WHERE ...

INSERT single & multiple documents

UPDATE any part of JSON document & use complex filter

DELETE

MERGE two sets of documents using traditional MERGE statement

EXPLAIN to understand the query plan

EXPLAIN SELECT \* FROM bucket WHERE ...

***Expressive, familiar, and feature-rich language for querying, transforming, and manipulating JSON data***



# N1QL (FEATURE-RICH)

---

Access to every part of JSON document

Functions (Date, Pattern, Array, Conditional, etc)

<https://developer.couchbase.com/documentation/server/current/n1ql/n1ql-language-reference/functions.html>

JOIN, NEST, UNNEST

Covering Index

Prepared Statements

USE KEYS, LIKE

***Expressive, familiar, and feature-rich language for querying, transforming, and manipulating JSON data***

# N1QL (Example)



*Dotted sub-document  
reference  
Names are CASE-SENSITIVE*

```
SELECT customers.id,  
        customers.NAME.lastname,  
        customers.NAME.firstname  
        Sum(orderline.amount)  
FROM   orders UNNEST orders.lineitems AS orderline  
        JOIN customers ON KEYS orders custid  
WHERE customers.state = 'NY'  
GROUP BY customers.id,  
          customers.NAME.lastname  
HAVING sum(orderline.amount) > 10000  
ORDER BY sum(orderline.amount) DESC
```

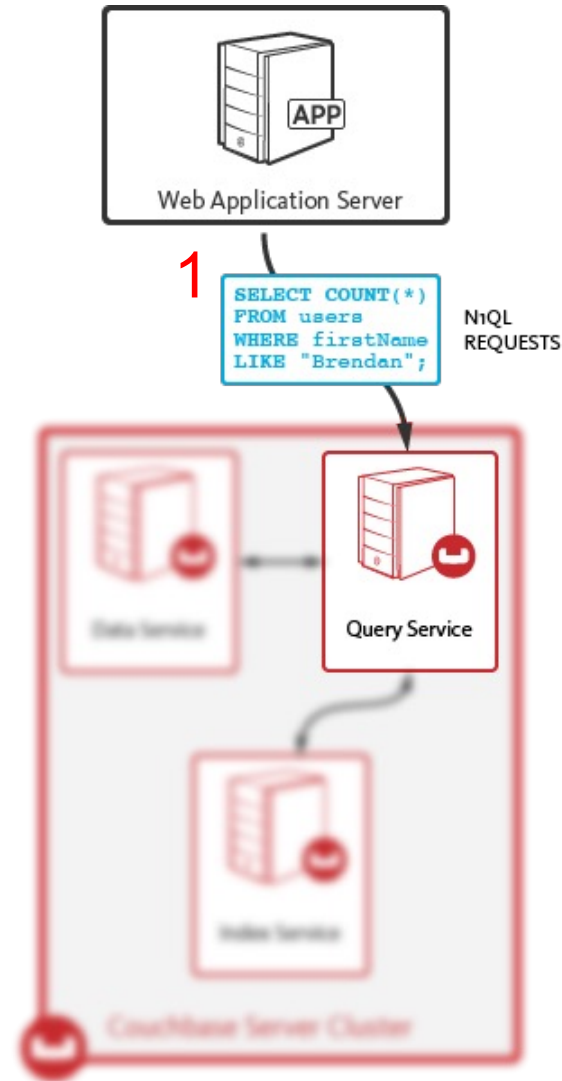
*UNNEST to flatten the arrays*

*JOINS with Document KEY of  
customers*



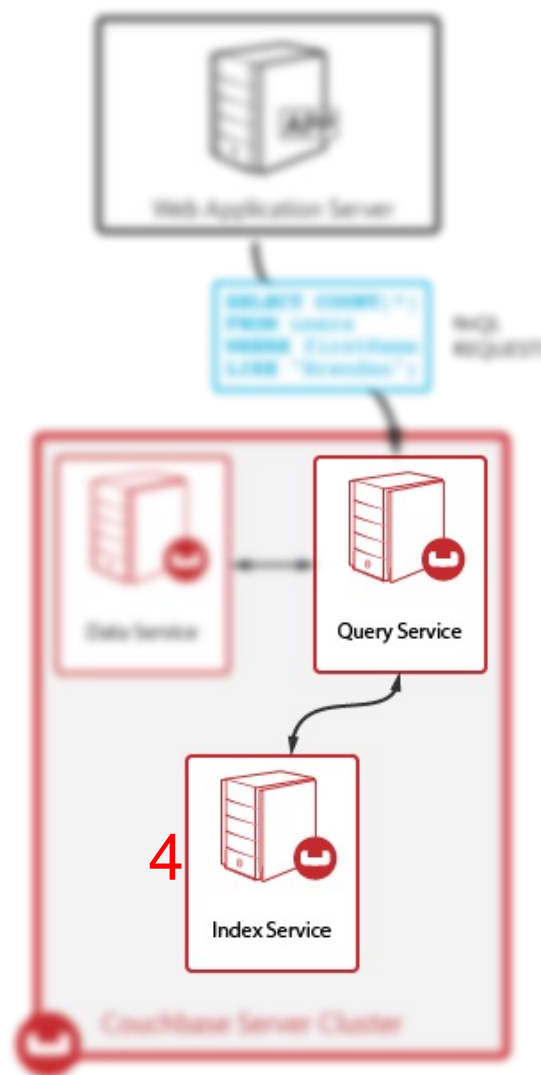
# N1QL Query : Execution

# Query Execution Flow



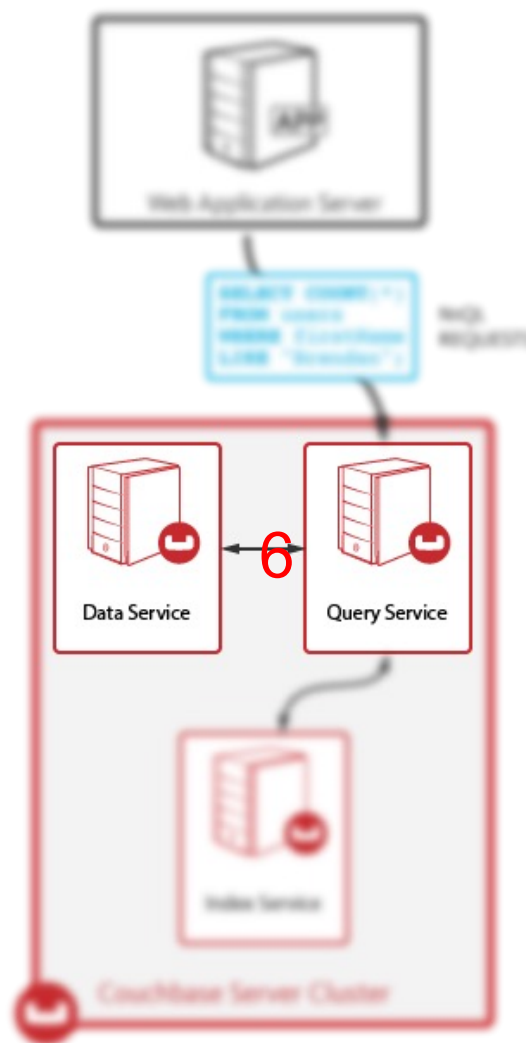
- 1. Application submits N1QL query**
- 2. Query is parsed, analyzed and plan is created**

# Query Execution Flow



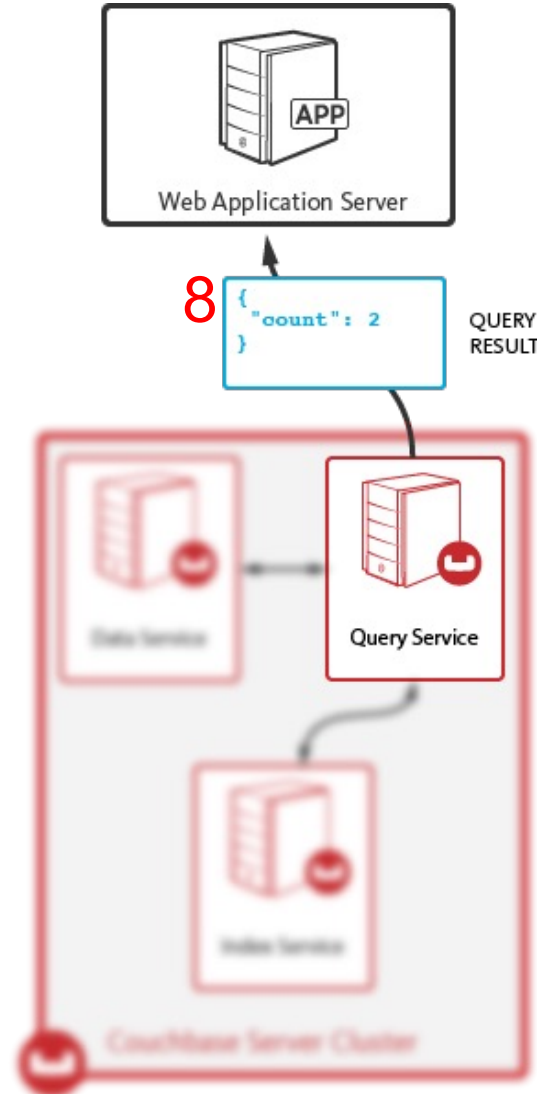
- 3. Query Service makes request to Index Service**
- 4. Index Service returns document keys and data**

# Query Execution Flow



- 5. If Covering Index, skip step 6**
- 6. If filtering is required, fetch documents from Data Service**

# Query Execution Flow

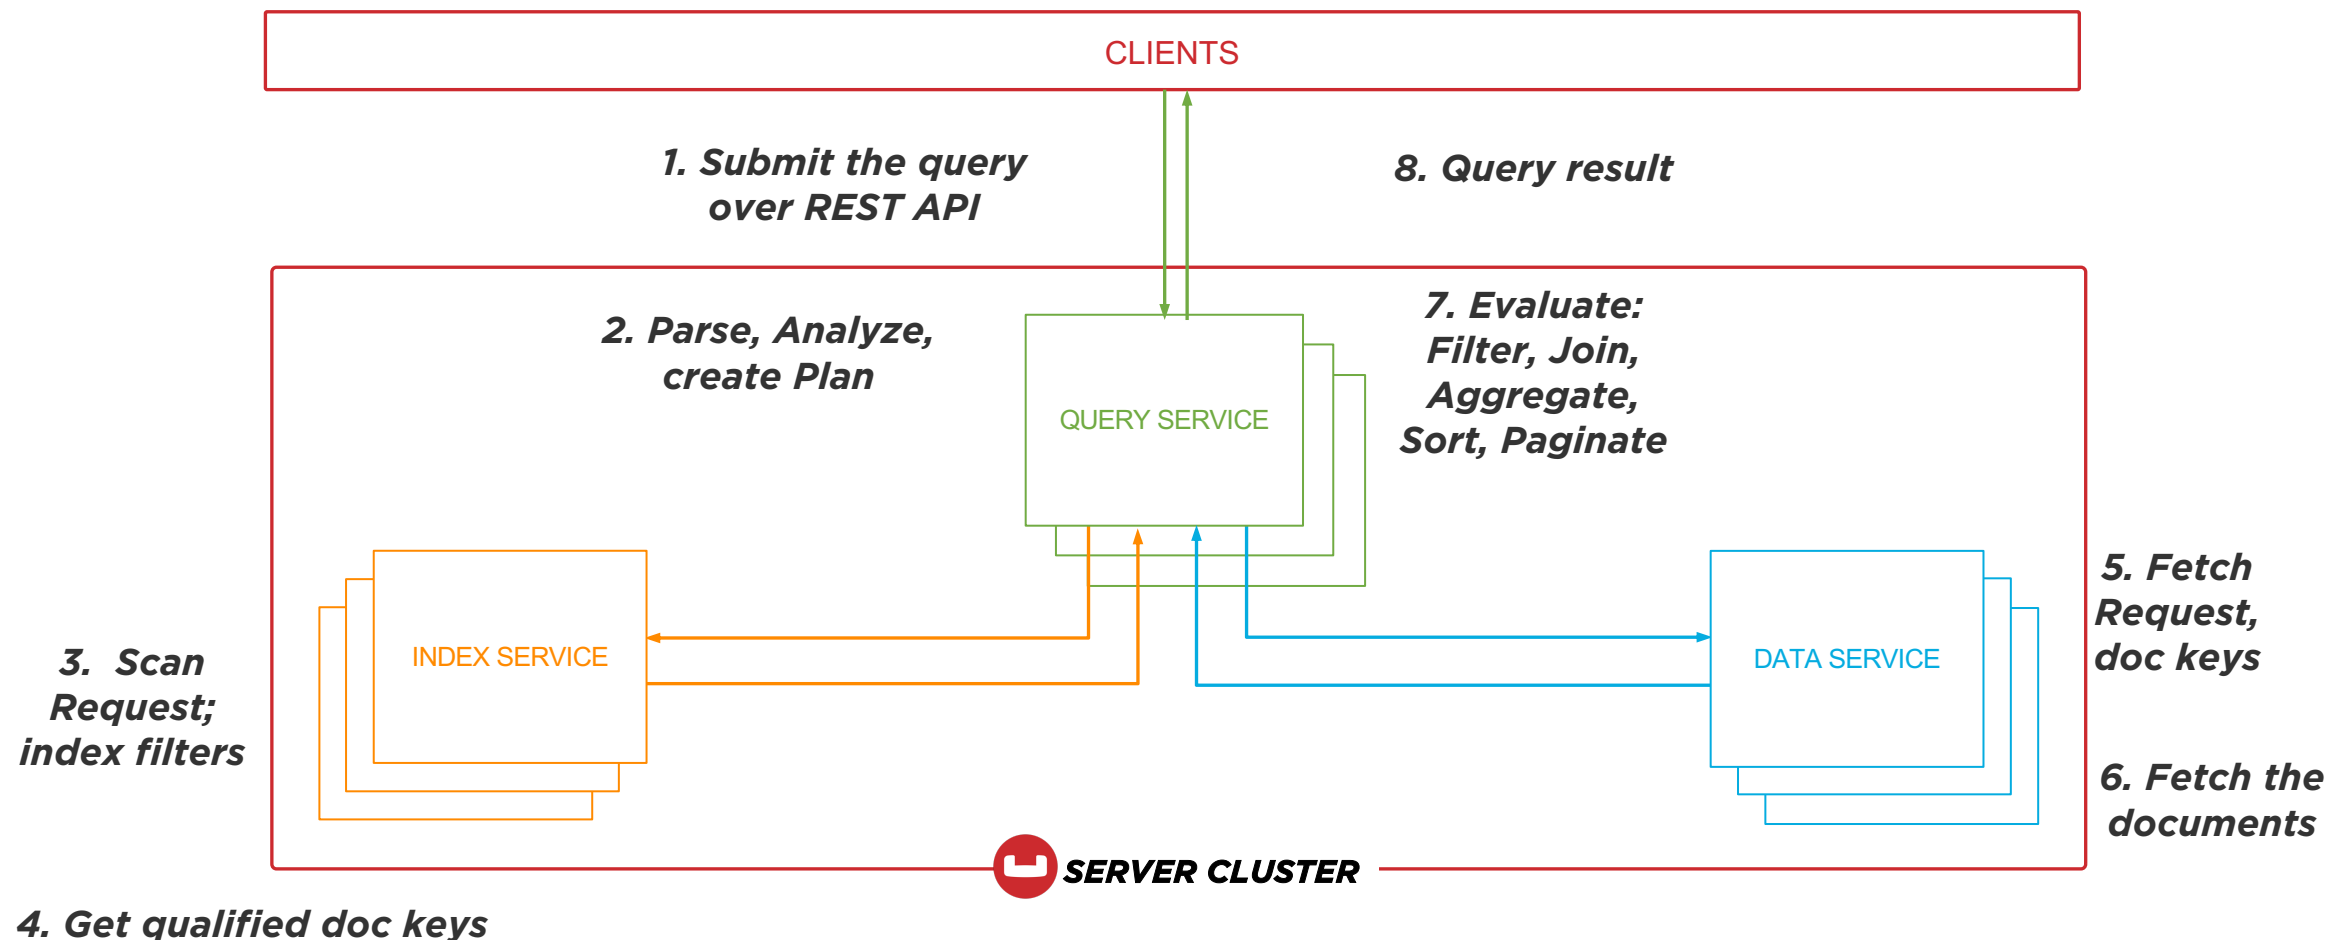


**7. Apply final logic  
(e.g. SORT, ORDER BY)**

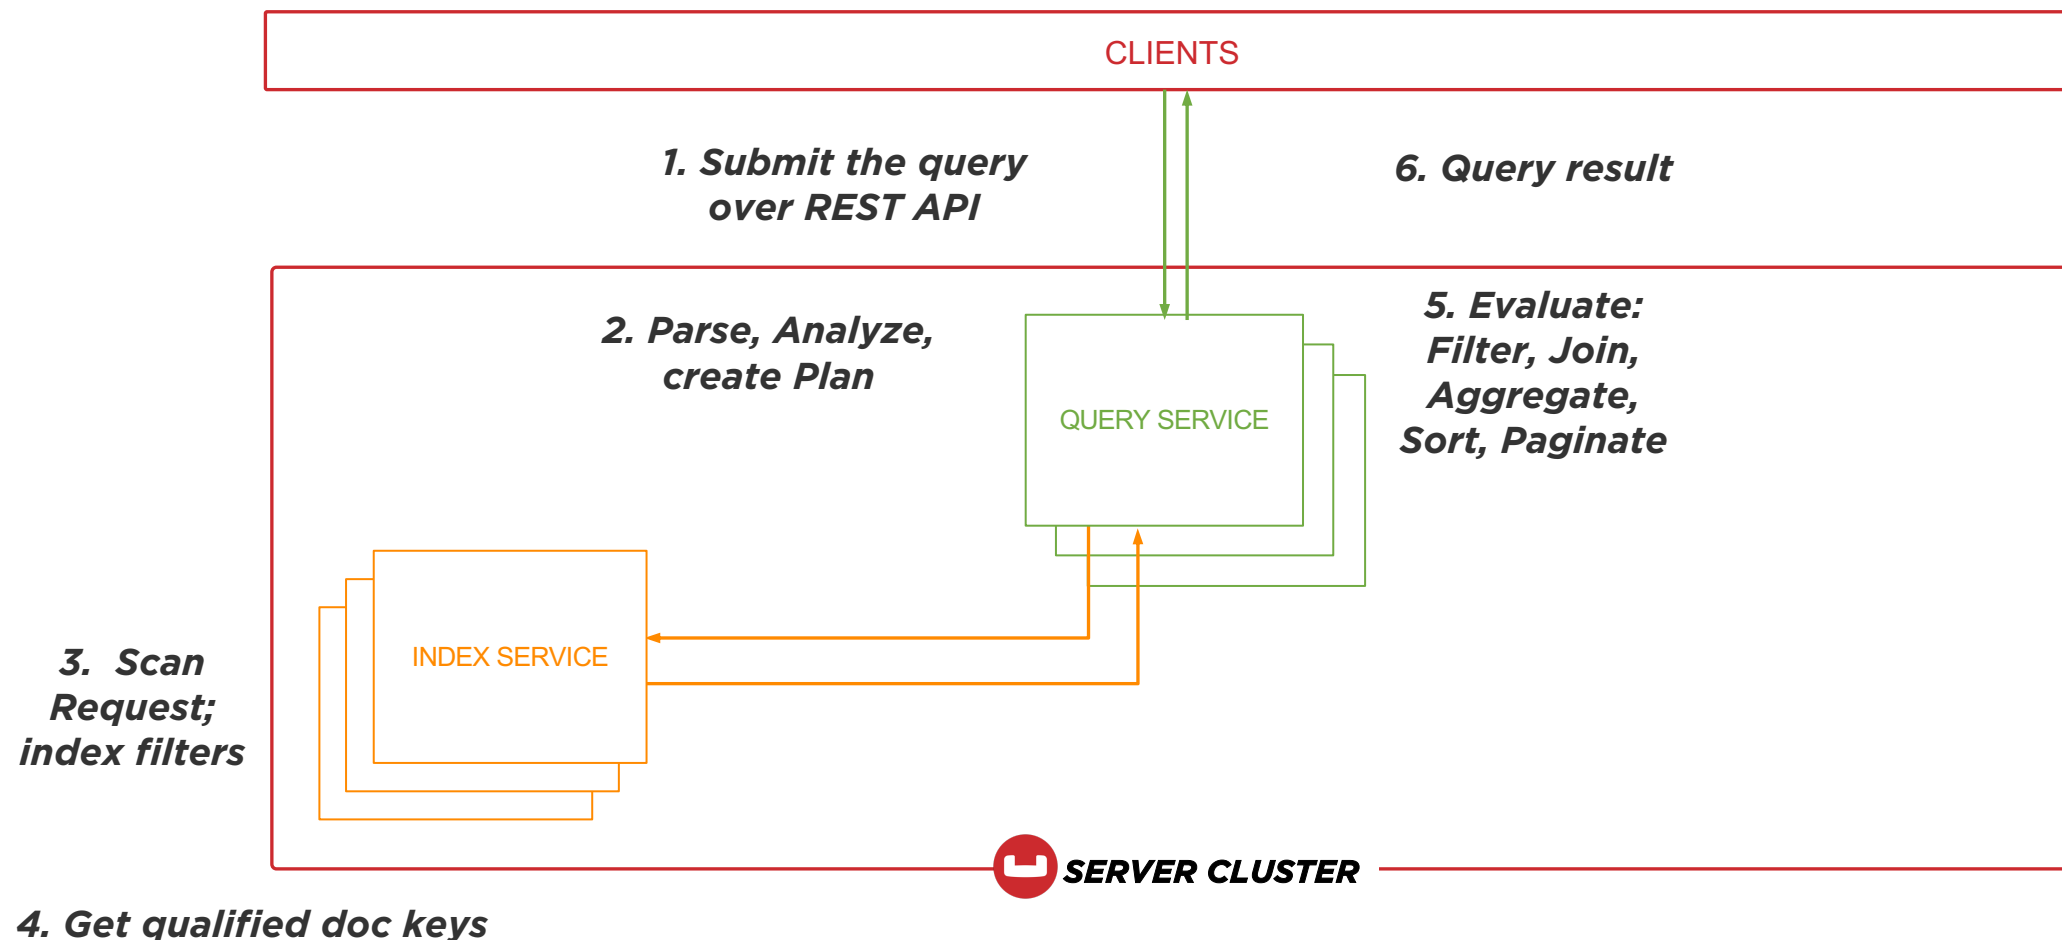
**8. Return formatted  
results to  
application**



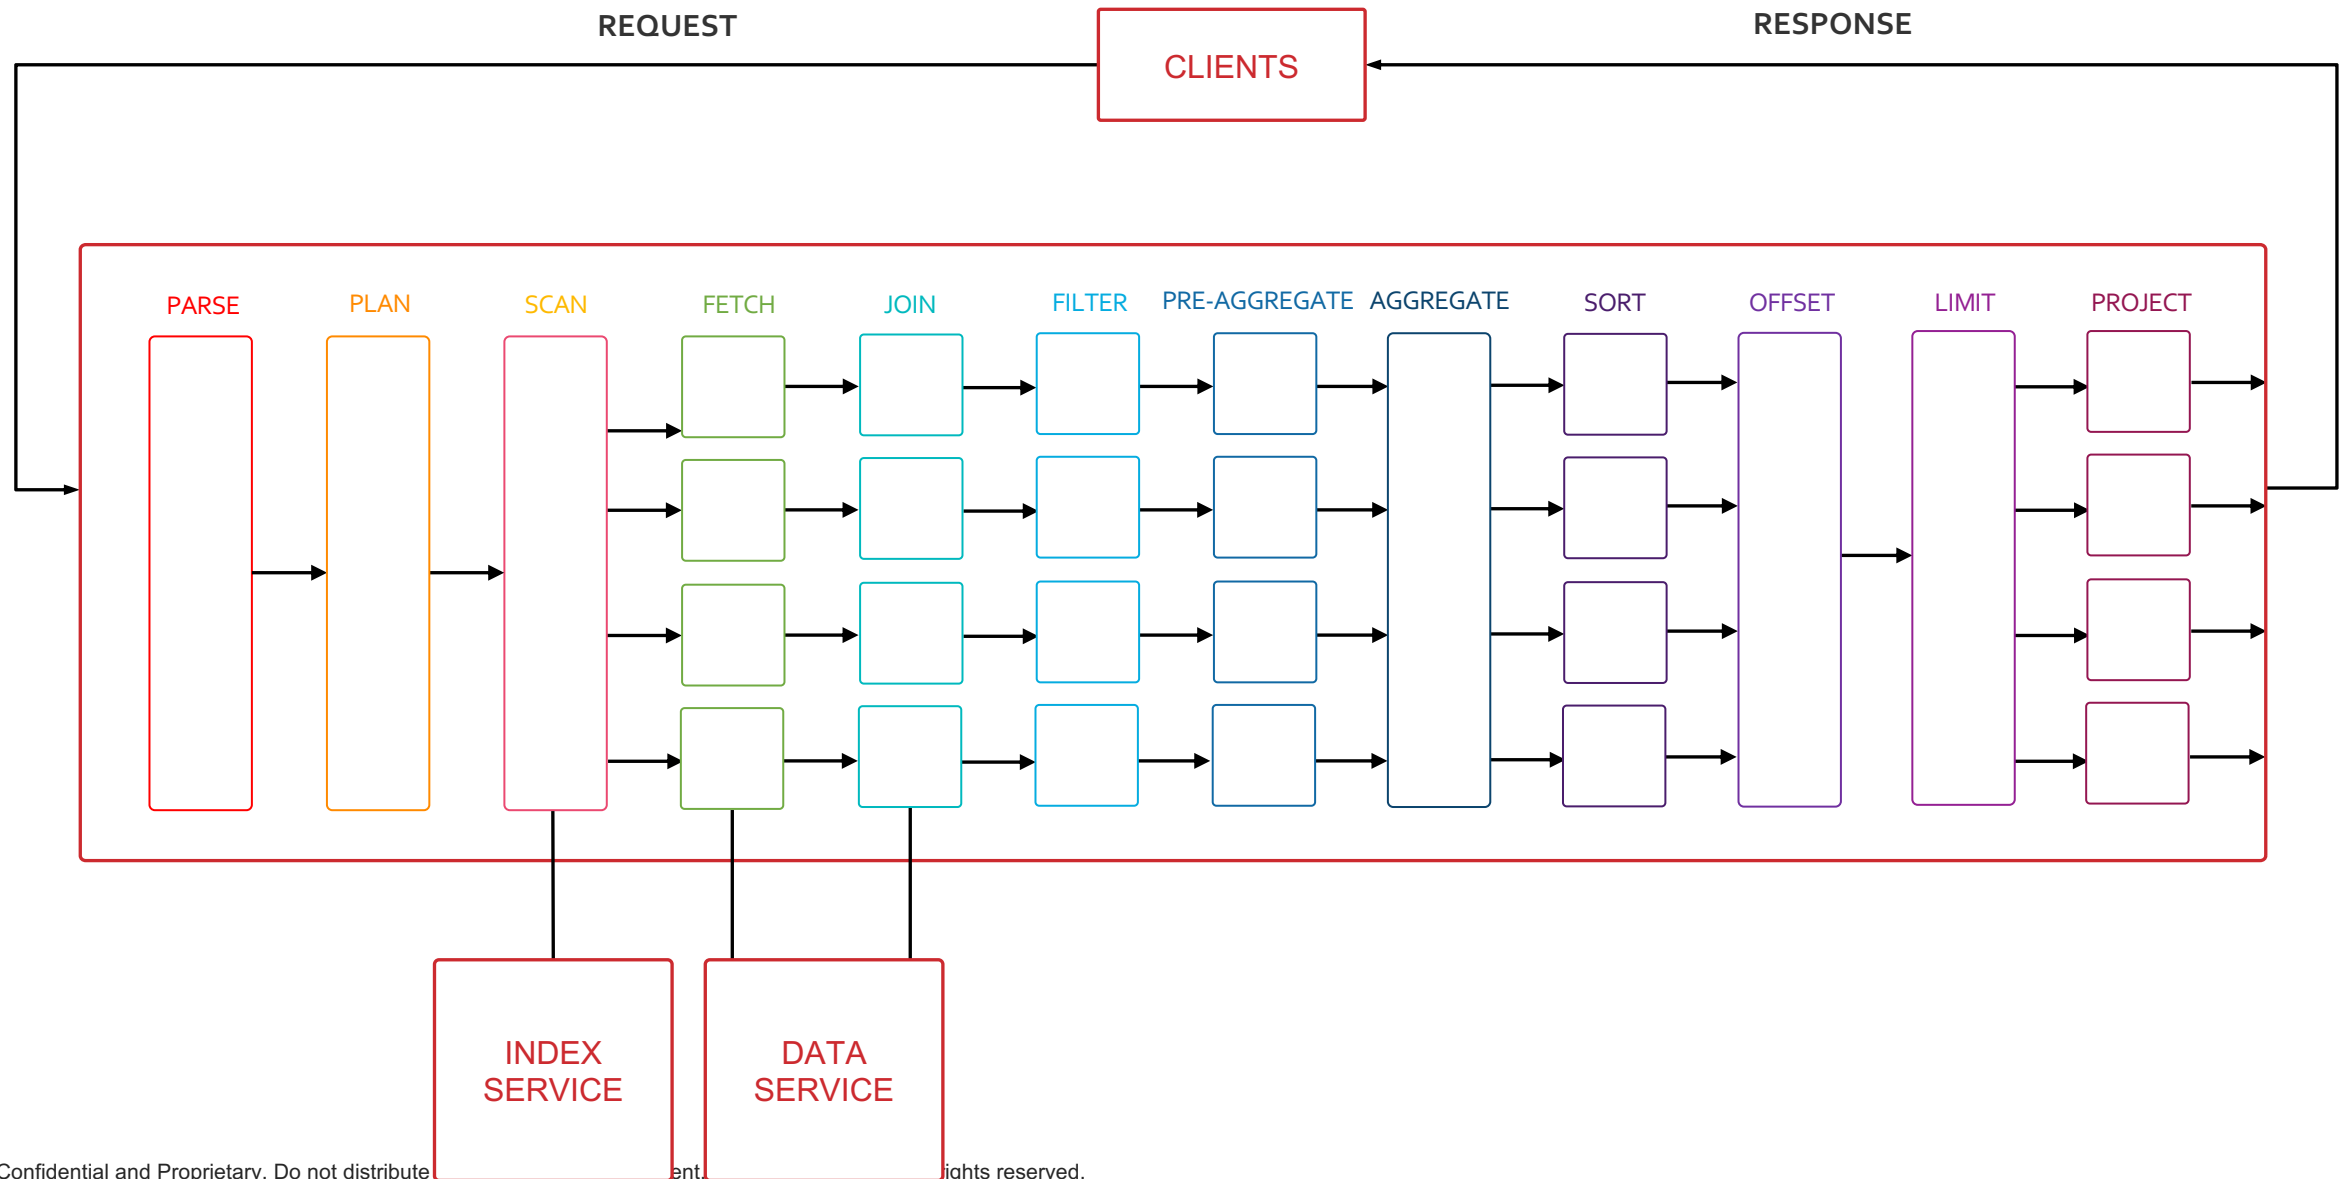
# Query Execution Flow (KV Fetch)



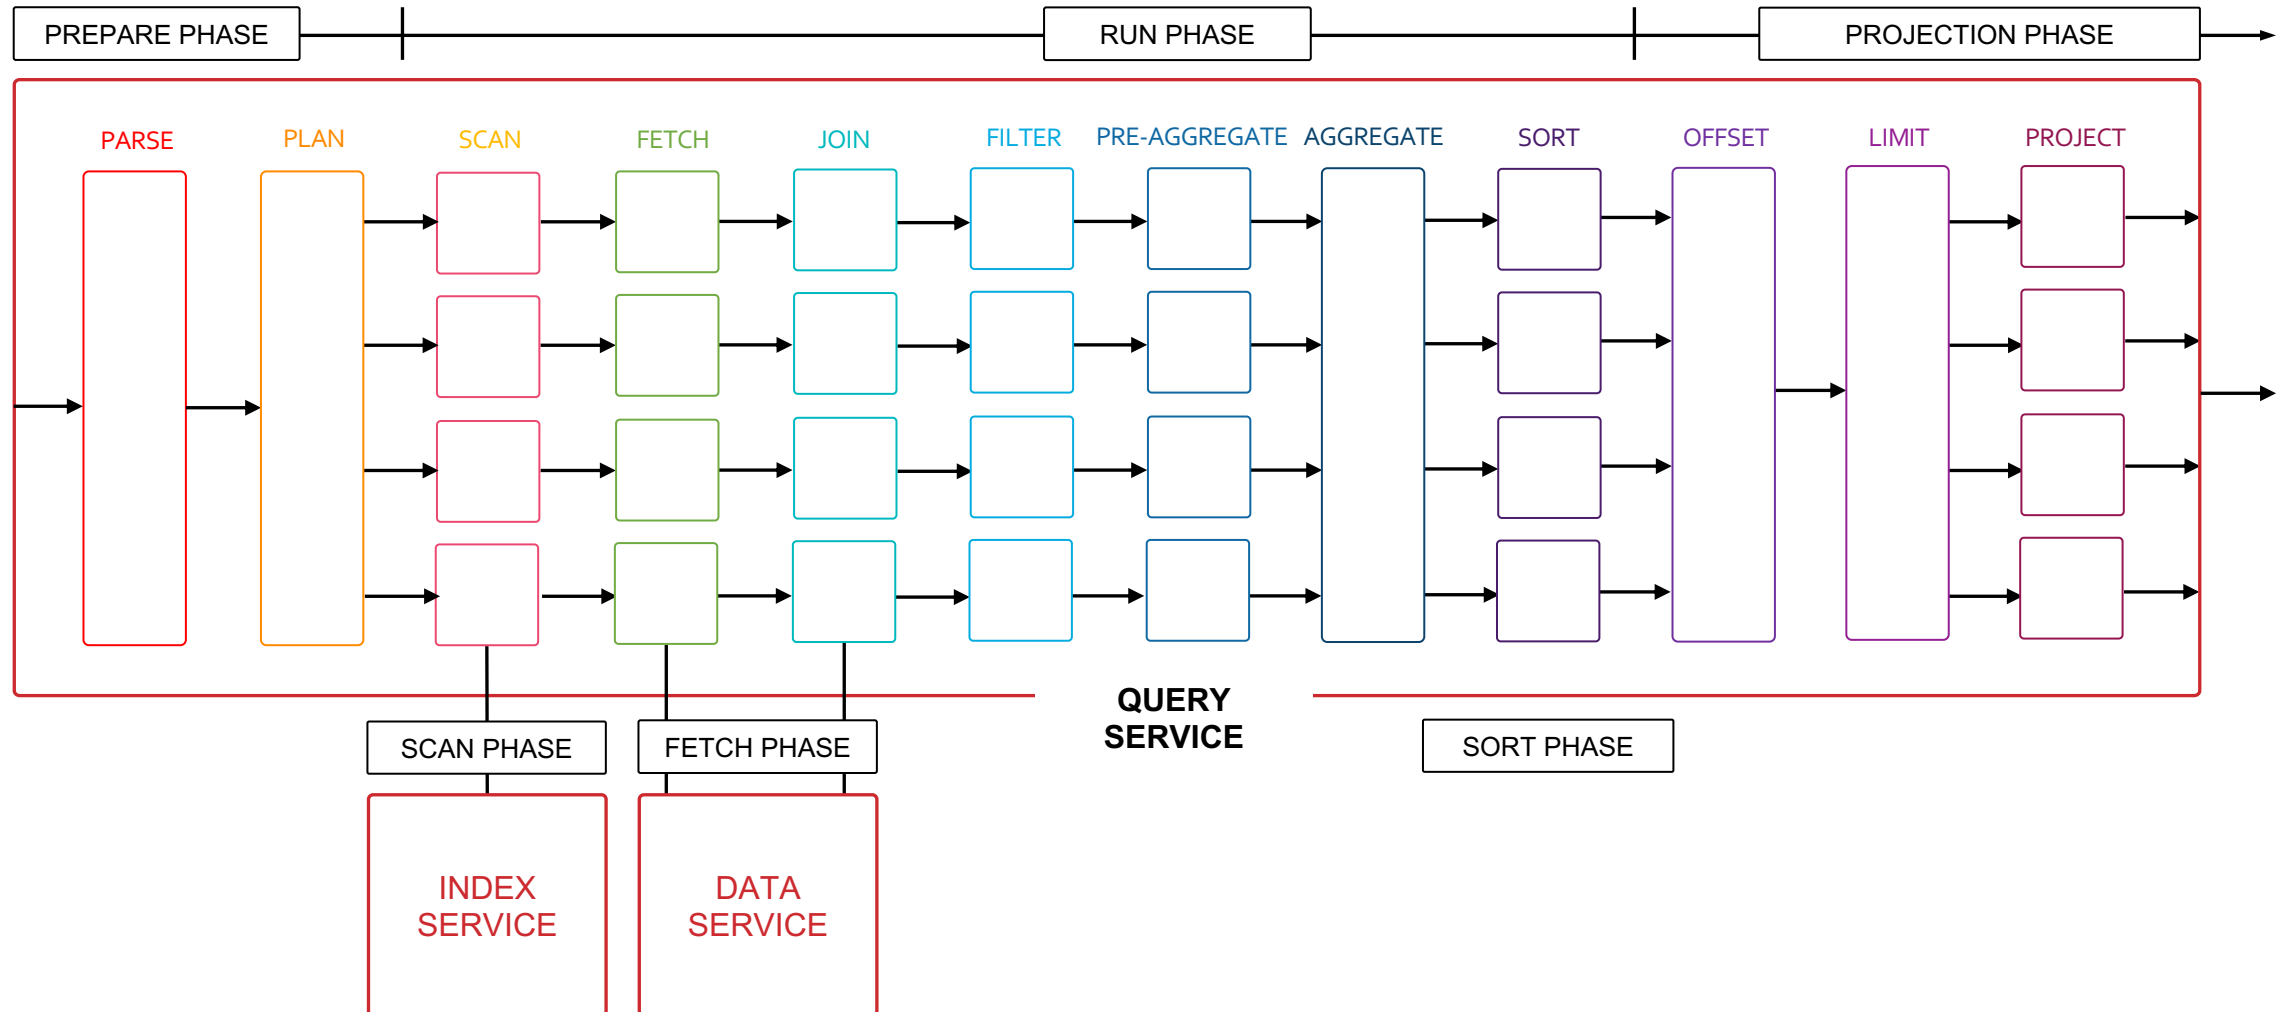
# Query Execution Flow (Covering Index)



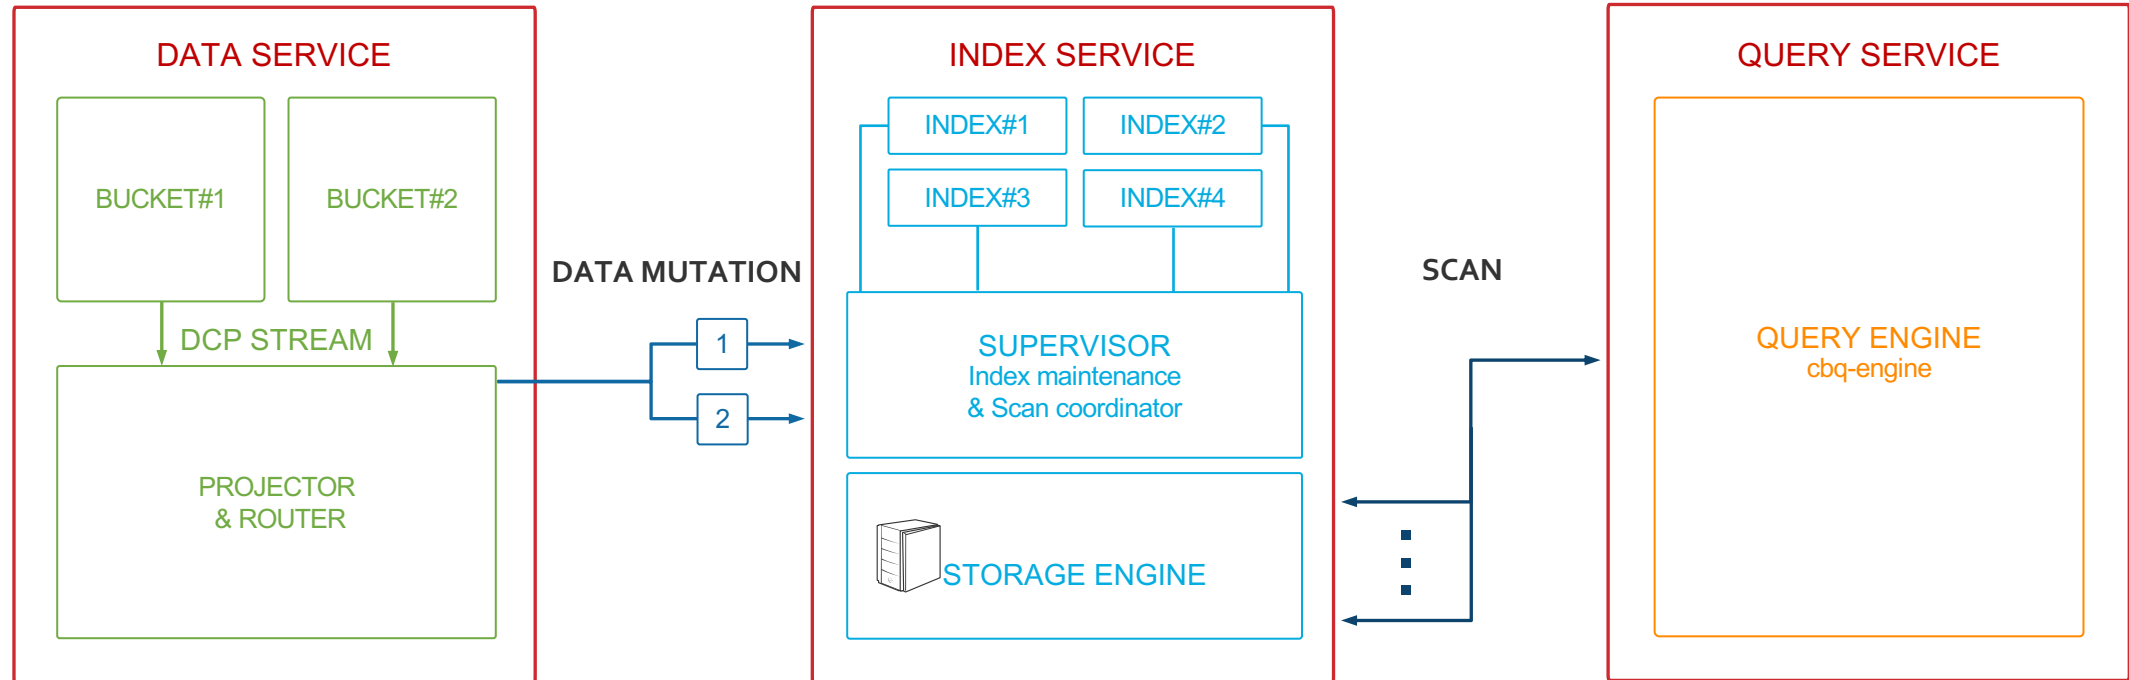
# Full Query Pipeline



# Inside a Query Service



# Couchbase Services



# Indexing



# Why do we need indexes?

---

***Indexes*** are used to  
***efficiently*** look up objects  
meeting ***user specified criteria***  
***without*** having to  
search ***every*** object  
in the collection



# Summary of Indexes

	Index Type	Description
1	<b>Primary</b>	Index on the document key on a whole bucket, to support full bucket scans
2	<b>Named</b>	Index with an assigned name, to allow multiple primary indexes in a cluster
3	<b>Secondary</b>	Index on a field (key-value pair) or document-key
4	<b>Composite</b>	Index on more than one field
5	<b>Functional</b>	Index on values resulting from a function or expression applied to a field
6	<b>Array</b>	Index on individual elements of array fields
7	<b>Partial</b>	Index on filtered subset of documents in a bucket
8	<b>Covering</b>	Describes any index which fully responds to a query without need for doc retrieval
9	<b>Duplicate</b>	Describes an indexing feature supporting load balancing, scale, and high availability
10	<b>Adaptive</b>	Describes an indexing feature providing an arrayed approach to generically indexing all or specified doc fields, to support increased query flexibility





# Secondary Indexes

---

Global Secondary Indexes (GSI) – Released in 4.0 (**Now called Legacy GSI / Deprecated in 5.0**)

- Lower query latency and high throughput

- Isolate from KV operations – multidimensional scaling

Memory Optimized Indexes (MOI) – new in 4.5

- Complete index is stored in memory.

- Supports much higher mutations and better performance

New Standard Global Secondary Index(GSI) – new in 5.0 (**AKA Plasma**)

- New storage engine



# N1QL (EXPLAIN)

*EXPLAIN statement when used before any N1QL statement, provides the execution plan for the statement*

***EXPLAIN SELECT log, type, runtime FROM logger ORDER BY runtime***

```
"results": [  
  {  
    "#operator": "Sequence",  
    "~children": [  
      {  
        "#operator": "Sequence",  
        "~children": [  
          {  
            "#operator": "PrimaryScan",  
            "index": "#primary",  
            "keyspace": "catalog",  
            "namespace": "default"  
          } ... and more...
```

# N1QL Language



# N1QL examples

---

```
SELECT d.C_ZIP, SUM(ORDLINE.OL_QUANTITY) AS TOTALQTY
FROM CUSTOMER d
      UNNEST ORDERS as CUSTORDERS
      UNNEST CUSTORDERS.ORDER_LINE AS ORDLINE
WHERE d.C_STATE = "NY"
GROUP BY d.C_ZIP
ORDER BY TOTALQTY DESC;
```

```
INSERT INTO CUSTOMER("PQR847", {"C_ID":4723, "Name":"Joe"});
```

```
UPDATE CUSTOMER c
SET c.STATE="CA", c.C_ZIP = 94501
WHERE c.ID = 4723;
```

# SELECT Statement

```
SELECT  customers.id,  
        customers.NAME.lastname,  
        customers.NAME.firstname  
        Sum(orderline.amount)  
  
FROM    orders UNNEST orders.lineitems AS orderline  
        JOIN customers ON KEYS orders.custid  
  
WHERE  customers.state = 'NY'  
  
GROUP BY customers.id,  
        customers.NAME.lastname  
  
HAVING  sum(orderline.amount) > 10000  
  
ORDER BY sum(orderline.amount) DESC
```

**Dotted sub-document reference  
Names are CASE-SENSITIVE**

**UNNEST to flatten the arrays**

**JOINS with Document KEY of  
customers**



# Composable SELECT statement

```
SELECT *
FROM
    (
        SELECT  a, b, c
        FROM    us_cust
        WHERE   x = 1
        ORDER BY x LIMIT 100 OFFSET 0
    UNION ALL
        SELECT  a, b, c
        FROM    canada_cust
        WHERE   y = 2
        ORDER BY x LIMIT 100 OFFSET 0    ) AS newtab
LEFT OUTER JOIN contacts
    ON KEYS newtab.c.contactid
ORDER BY a, b, c
LIMIT 10 OFFSET 100
```



# SELECT Statement Highlights

---

## Querying across relationships

JOINS

Subqueries

## Aggregation

MIN, MAX

SUM, COUNT, AVG, ARRAY\_AGG [ DISTINCT ]

## Combining result sets using set operators

UNION, UNION ALL, INTERSECT, EXCEPT



# N1QL Query Operators [ 1 of 2 ]

---

## USE KEYS ...

Direct primary key lookup bypassing index scans

Ideal for hash-distributed datastore

Available in SELECT, UPDATE, DELETE

## JOIN ... ON KEYS ...

Nested loop JOIN using key relationships

Ideal for hash-distributed datastore

Current implementation supports INNER and LEFT OUTER joins





# N1QL Query Operators [ 2 of 2 ]

---

## NEST

Special JOIN that embeds external child documents under their parent  
Ideal for JSON encapsulation

## UNNEST

Flattening JOIN that surfaces nested objects as top-level documents  
Ideal for decomposing JSON hierarchies

JOIN, NEST, and UNNEST can be chained in any combination



# N1QL Expressions for JSON

<b><i>Ranging over collections</i></b>	<ul style="list-style-type: none"><li>• <b><i>WHERE ANY c IN children SATISFIES c.age &gt; 10 END</i></b></li><li>• <b><i>WHERE EVERY r IN ratings SATISFIES r &gt; 3 END</i></b></li></ul>
<b><i>Mapping with filtering</i></b>	<ul style="list-style-type: none"><li>• <b><i>ARRAY c.name FOR c IN children WHEN c.age &gt; 10 END</i></b></li></ul>
<b><i>Deep traversal, SET, and UNSET</i></b>	<ul style="list-style-type: none"><li>• <b><i>WHERE ANY node WITHIN request SATISFIES node.type = "xyz" END</i></b></li><li>• <b><i>UPDATE doc UNSET c.field1 FOR c WITHIN doc END</i></b></li></ul>
<b><i>Dynamic Construction</i></b>	<ul style="list-style-type: none"><li>• <b><i>SELECT { "a": expr1, "b": expr2 } AS obj1, name FROM ... // Dynamic object</i></b></li><li>• <b><i>SELECT [ a, b ] FROM ... // Dynamic array</i></b></li></ul>
<b><i>Nested traversal</i></b>	<ul style="list-style-type: none"><li>• <b><i>SELECT x.y.z, a[0] FROM a.b.c ...</i></b></li></ul>
<b><i>IS [ NOT ] MISSING</i></b>	<ul style="list-style-type: none"><li>• <b><i>WHERE name IS MISSING</i></b></li></ul>



# N1QL Data Types from JSON

---

***N1QL supports all JSON data types***

***Numbers***

***Strings***

***Booleans***

***Null***

***Arrays***

***Objects***



# N1QL Data Type Handling

---

## Non-JSON data types

MISSING

Binary

## Data type handling

Date functions for string and numeric encodings

Total ordering across all data types

Well defined semantics for ORDER BY and comparison operators

Defined expression semantics for all input data types

No type mismatch errors



# Data Modification Statements

---

UPDATE ... SET ... WHERE ...

DELETE FROM ... WHERE ...

INSERT INTO ... ( KEY, VALUE ) VALUES ...

INSERT INTO ... ( KEY ..., VALUE ... ) SELECT ...

MERGE INTO ... USING ... ON ...

WHEN [ NOT ] MATCHED THEN ...

Note: Couchbase provides per-document atomicity.

# Data Modification Statements

***JSON literals can be used in any expression***



***INSERT INTO ORDERS (KEY, VALUE)***

***VALUES ("1.ABC.X382", {"O\_ID":482, "O\_D\_ID":3, "O\_W\_ID":4});***

***UPDATE ORDERS***

***SET O\_CARRIER\_ID = "ABC987"***

***WHERE O\_ID = 482 AND O\_D\_ID = 3 AND O\_W\_ID = 4***

***DELETE FROM NEW\_ORDER***

***WHERE NO\_D\_ID = 291 AND***

***NO\_W\_ID = 3482 AND***

***NO\_O\_ID = 2483***



# Index Statements

---

***CREATE INDEX ON ...***

***DROP INDEX ...***

***EXPLAIN ...***

***Highlights***

***Functional indexes***

***on any data expression***

***Partial indexes***



# Index Overview: Secondary Index

Document key: "customer534"

```
"customer": {  
  "ccInfo": {  
    "cardExpiry": "2015-11-11",  
    "cardNumber": "1212-232-1234",  
    "cardType": "americanexpress",  
  },  
  "customerId": "customer534",  
  "dateAdded": "2014-04-06",  
  "dateLastActive": "2014-05-02",  
  "emailAddress": "iles@kertz.name",  
  "firstName": "Mckayla",  
  "lastName": "Brown",  
  "phoneNumber": "1-533-290-6403",  
  "postalCode": "92341",  
  "state": "VT",  
  "type": "customer"  
}
```

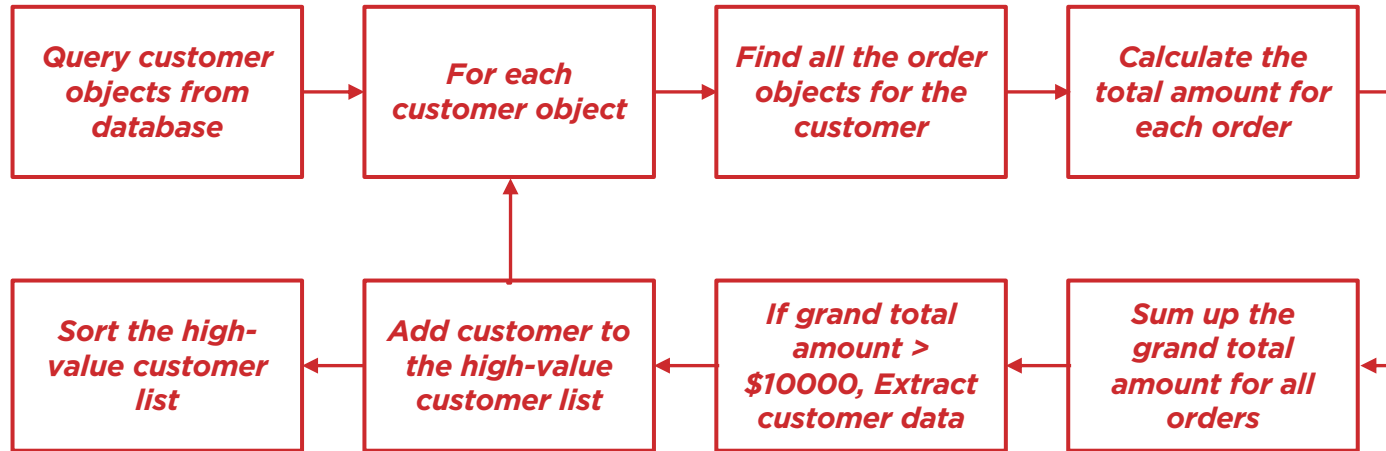
- ***Secondary Index can be created on any combination of expressions.***
  - ***CREATE INDEX idx\_cust\_cardnum customer(ccInfo.cardNumber, postalcode)***
- ***Useful in speeding up the queries.***
- ***Need to have matching indexes with matching key-ordering***
  - ***(ccInfo.cardExpiry, postalCode)***
  - ***(type, state, lastName firstName)***





# Find High-Value Customers with Orders > \$10000

LOOPING OVER MILLIONS OF CUSTOMERS IN APPLICATION!!!



- **Complex codes and logic**
- **Inefficient processing on client side**

**API QUERY**

**VS.**

**SQL for JSON**

```
SELECT Customers.ID, Customers.Name, SUM(OrderLine.Amount)
FROM Orders UNNEST Orders.LineItems AS OrderLine
JOIN Customers ON KEYS Orders.CustID
GROUP BY Customers.ID, Customers.Name
HAVING SUM(OrderLine.Amount) > 10000
ORDER BY SUM(OrderLine.Amount) DESC
```

- **Proven and expressive query language**
- **Leverage SQL skills and ecosystem**
- **Extended for JSON**

# Summary: SQL & N1QL



<b>Query Features</b>	<b>SQL</b>	<b>N1QL</b>
<b>Statements</b>	<ul style="list-style-type: none"><li>▪ <b><i>SELECT, INSERT, UPDATE, DELETE, MERGE</i></b></li></ul>	<ul style="list-style-type: none"><li>▪ <b><i>SELECT, INSERT, UPDATE, DELETE, MERGE</i></b></li></ul>
<b>Query Operations</b>	<ul style="list-style-type: none"><li>▪ <b><i>Select, Join, Project, Subqueries</i></b></li><li>▪ <b><i>Strict Schema</i></b></li><li>▪ <b><i>Strict Type checking</i></b></li></ul>	<ul style="list-style-type: none"><li>▪ <b><i>Select, Join, Project, Subqueries</i></b></li><li>✓ <b><i>Nest &amp; Unnest</i></b></li><li>✓ <b><i>Look Ma! No Type Mismatch Errors!</i></b></li><li>▪ <b><i>JSON keys act as columns</i></b></li></ul>
<b>Schema</b>	<ul style="list-style-type: none"><li>▪ <b><i>Predetermined Columns</i></b></li></ul>	<ul style="list-style-type: none"><li>✓ <b><i>Fully addressable JSON</i></b></li><li>✓ <b><i>Flexible document structure</i></b></li></ul>
<b>Data Types</b>	<ul style="list-style-type: none"><li>▪ <b><i>SQL Data types</i></b></li><li>▪ <b><i>Conversion Functions</i></b></li></ul>	<ul style="list-style-type: none"><li>▪ <b><i>JSON Data types</i></b></li><li>▪ <b><i>Conversion Functions</i></b></li></ul>
<b>Query Processing</b>	<ul style="list-style-type: none"><li>▪ <b><i>INPUT: Sets of Tuples</i></b></li><li>▪ <b><i>OUTPUT: Set of Tuples</i></b></li></ul>	<ul style="list-style-type: none"><li>▪ <b><i>INPUT: Sets of JSON</i></b></li><li>▪ <b><i>OUTPUT: Set of JSON</i></b></li></ul>



# Using N1QL: Index Scan

```
CREATE INDEX `idx_id` ON `travel-sample`(`id`);
```

```
EXPLAIN SELECT meta().id FROM `travel-sample` WHERE id = 10;
```

```
EXPLAIN SELECT meta().id FROM `travel-sample` WHERE id >=10 AND id < 25;
```

```
EXPLAIN SELECT meta().id FROM `travel-sample` WHERE id IN [10, 20];
```

<b><i>Predicate</i></b>	<b><i>Span Low</i></b>	<b><i>Span High</i></b>	<b><i>Inclusion</i></b>
<b><i>id = 10</i></b>	<b><i>10</i></b>	<b><i>10</i></b>	<b><i>3 (Both)</i></b>
<b><i>id &gt;= 10</i></b>	<b><i>10</i></b>	<b><i>None</i></b>	<b><i>0 (Neither)</i></b>
<b><i>id &lt;= 10</i></b>	<b><i>Null</i></b>	<b><i>10</i></b>	<b><i>2 (High)</i></b>



## Using N1QL: LIKE

```
select SUFFIXES("Baltimore Washington Intl")  
create index suffixes_airport_name on `travel-sample`  
(DISTINCT ARRAY array_element FOR array_element IN  
SUFFIXES(LOWER(airportname)) END)  
WHERE type = "airport";
```

```
SELECT airportname FROM `travel-  
sample`  
WHERE airportname LIKE  
"%Washington%"  
AND type = "airport"
```


# Using N1QL: JOIN




*Inner, Left, Right, Outer, Exclude,*

***SELECT <select\_list>  
FROM bucket A  
LEFT JOIN bucket B  
ON KEYS <keys-  
clause(A)>***


## N1QL JOINS




```
SELECT <select_list>  
FROM Table_A A  
LEFT JOIN Table_B B  
ON KEYS <keys-clause(A)>
```




```
SELECT <select_list>  
FROM Table_A A  
INNER JOIN Table_B B  
ON KEYS <keys-clause(A)>
```




```
SELECT <select_list>  
FROM Table_B B  
LEFT JOIN Table_A A  
ON KEYS <keys-clause(B)>
```




```
SELECT <select_list>  
FROM Table_B B  
LEFT JOIN Table_A A  
ON KEYS <keys-clause(B)>  
WHERE META(A).id IS MISSING  
UNION ALL  
SELECT <select_list>  
FROM Table_A A  
LEFT JOIN Table_B B  
ON KEYS <keys-clause(A)>  
WHERE META(B).id IS MISSING
```



```
SELECT <select_list>  
FROM Table_B B  
LEFT JOIN Table_A A  
ON KEYS <keys-clause(B)>  
WHERE META(B).id IS MISSING
```



```
SELECT <select_list>  
FROM Table_A A  
LEFT JOIN Table_B B  
ON KEYS <keys-clause(A)>  
WHERE META(B).id IS MISSING
```



```
SELECT <select_list>  
FROM Table_B B  
LEFT JOIN Table_A A  
ON KEYS <keys-clause(B)>  
UNION ALL  
SELECT <select_list>  
FROM Table_A A  
LEFT JOIN Table_B B  
ON KEYS <keys-clause(A)>  
WHERE META(B).id IS MISSING
```

@atom\_yang, 2017

# Query Consistency



***T3: issue query on (k1,v1)***

***...***

***T2: do other business logic computation***

***...***

***T1: insert (k1, v1)***

## ***Strict Request-Time Consistency***

***(request\_plus)***

***Query execution is delayed until all indexes process mutations up to T3***

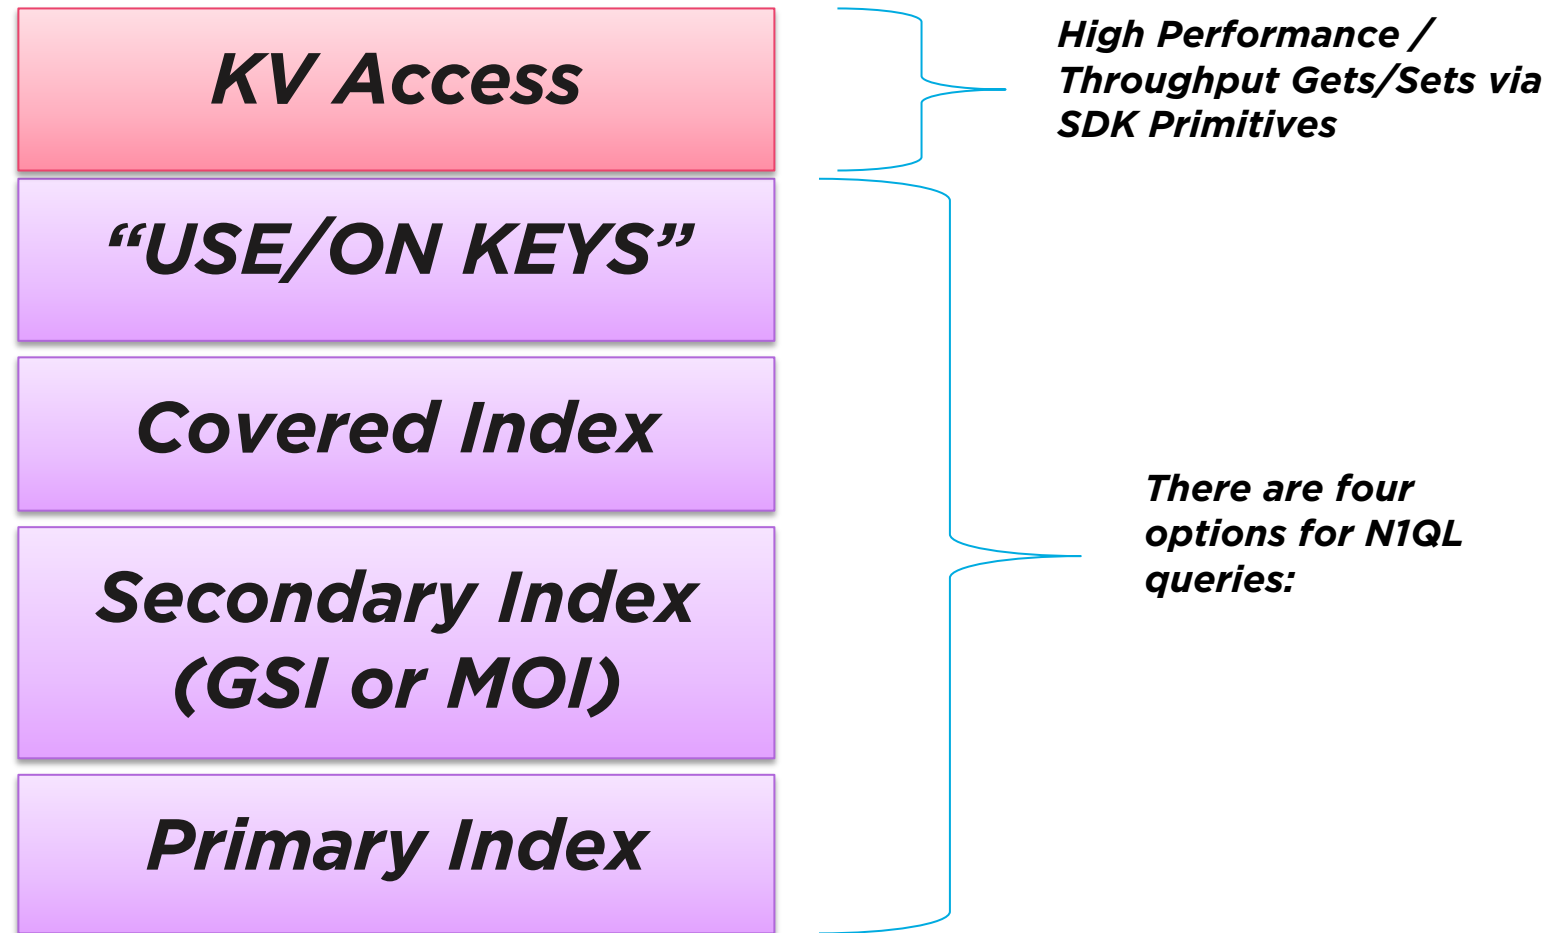
## ***RYOW Consistency***

***(at\_plus)***

***Query execution is delayed until all indexes process mutations up to T1***

# Query Tuning

# Optimizing Query Speed







# Query Tuning

---

- Rule based optimization
- Optimal indexes for optimal performance
- EXPLAIN to understand and tune
- Examine configuration for Data, Indexer, Query



# Query Tuning List

---

- Query should have predicates to avoid primary index scan.
- Explore all index options (composite secondary index, partial composite index, covered partial composite index ,...)
- Include as many predicate keys as possible in leading index keys
  - Query processing is bit more efficient when there is equality predicates on leading index keys
- Explore avoiding Intersect scan. If required provide hint with USE INDEX
- For ANY predicate clause, use ARRAY index keys.
- Explore using index key order and pushing limit, offset to indexer
- Rewrite query if required
- Use array fetch when possible
- Execute query and look the monitoring stats for each phase of query (ex: system:completed\_requests) and tune it.
- Check the final query plan through the explain
- Set pretty=false in Query Service or query parameter
- Increase parallel processing via max\_parallelism
- Fetching large set of data for non covered index increase pipeline-cap, pipeline-batch in Query-Service
- Duplicate indexes and memory optimized indexes
- Add more Query nodes or Indexer nodes



# Pattern “USE KEYS”

***USE KEYS provides facility Query Service to access Data Service directly***

***No index required***

***Scales independent of bucket size***

```
SELECT b.destinationairport, b.sourceairport,  
(SELECT c.name  
FROM `travel-sample` c  
USE KEYS b.airlineid  
LEFT UNNEST c.name) as theAirline  
FROM `travel-sample` b  
USE KEYS "route_5966"
```

# Pattern “Covering Index”

Index selection for a query solely depends on the query predicates

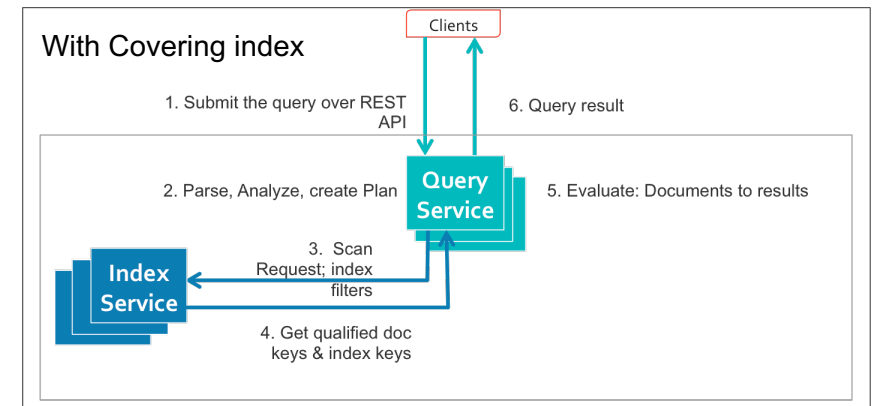
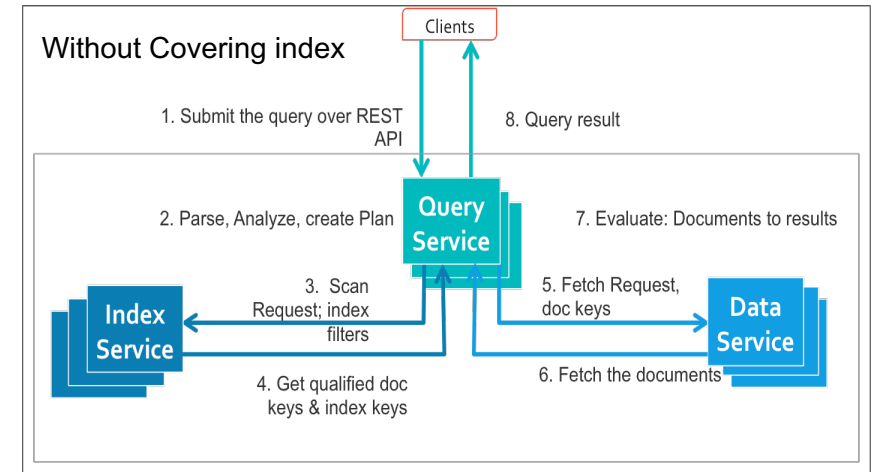
Index keys cover predicates and all attribute references

Avoids fetching the whole document

Performance

```
CREATE INDEX ts_airline ON `travel-sample`(name, id)
WHERE type = "airline";
```

```
SELECT name, id
FROM `travel-sample`
WHERE type = "airline" AND name = "United Airlines";
```





# Pattern “Array Index and Query”

## Ad Targeting System

Raw Logs

User Interaction Log

User Profile (Ref Doc)

Add an Index (you can do this at any time!)

```
CREATE INDEX iArray ON default(distinct (array v for v in
    interactions end))
```

```
SELECT meta().id from default WHERE ANY v IN interactions
    SATISFIES v.last_interaction='Wait' END;
```

```
Hailie20.Doyle34@hotmail.com::interactions
[
  {
    "default": {
      "interactions": [
        {
          "interaction_id": "04ec514a-d6c2-42b7-962a-
bf8976677e79",
          "interaction_type": "wait"
        }
      ]
    },
    {
      "interaction_id": "04ec514a-d6c2-42b7-962a-
ac9467788e88",
      "interaction_type": "email"
    }
  }
  ...
]
```



# Monitoring: Active requests

---

List / Delete requests currently being run by the query service

Through N1QL

```
SELECT * FROM system:active_requests
```

```
DELETE FROM system:active_requests WHERE...
```

Through REST

```
GET http://localhost:8093/admin/active\_requests
```

```
GET http://localhost:8093/admin/active\_requests/<request\_id>
```

```
DELETE http://localhost:8093/admin/active\_requests/<request\_id>
```



# Monitoring: Prepared statements

---

List / Delete requests prepared on the query node

Through N1QL

```
SELECT * FROM system:prepareds
```

```
DELETE FROM system:prepareds WHERE...
```

Through REST

```
GET http://localhost:8093/admin/prepareds
```

```
GET http://localhost:8093/admin/prepareds/<request_id>
```

```
DELETE http://localhost:8093/admin/prepareds/<request_id>
```



# Monitoring: Completed requests

---

List / Delete completed requests deemed to be of high cost

Through N1QL

```
SELECT * FROM system:completed_requests
```

```
DELETE FROM system:completed_requests where...
```

Through REST

```
GET http://localhost:8093/admin/completed\_requests
```

```
GET http://localhost:8093/admin/completed\_requests/<request\_id>
```

```
DELETE http://localhost:8093/admin/completed\_requests/<request\_id>
```

Provide an overall health picture of the query service

```
Using REST: GET http://localhost:8093/admin/vitals
```





# Thank You