



# Architecture and Administration Basics

Java Labs

1

# Installation & Configuration SDK

# Documentation & Examples

---

Open the documentation for `libcouchbase`!

- <https://developer.couchbase.com/documentation/server/current/sdk/java/sample-app-backend.html>
- <https://developer.couchbase.com/documentation/server/current/sdk/java/start-using-sdk.html>
- <http://docs.couchbase.com/sdk-api/couchbase-java-client-2.7.2/>
- Check the Sample App backend  
<https://docs.couchbase.com/java-sdk/2.7/sample-application.html>

# Including the SDK

```
// Gradle  
dependencies {  
    ...  
    compile group: 'com.couchbase.client', name: 'java-client', version:'2.5.4'  
    ...  
}
```

```
// Maven  
<dependencies>  
    <dependency>  
        <groupId>com.couchbase.client</groupId>  
        <artifactId>java-client</artifactId>  
        <version>2.7.3</version>  
    </dependency>  
</dependencies>
```

**TIP.** Check last version available here:

<https://developer.couchbase.com/documentation/server/current/sdk/java/start-using-sdk.html>

# Lab 1: Preparing for the Lab

- Clone the source repository in your home folder.  
`git clone https://github.com/vgasiusas/cb-workshop-2d.git`
- Make sure Couchbase Server is running on your Machine and you have travel-sample bucket
- Create a user for the application access
  - Go the Security Tab in Couchbase
  - Create a new User.
  - Username = "travel"
  - Password = "couchbase"
  - Role: Application Access on 'travel-sample'

Roles

<input type="checkbox"/> Query System Catalog
<input type="checkbox"/> Analytics Reader
<input type="checkbox"/> All Buckets (*)
<input type="checkbox"/> travel-sample
<input type="checkbox"/> Bucket Admin
<input checked="" type="checkbox"/> Application Access <span style="color: green;">✓</span>
<input type="checkbox"/> XDCR Inbound
<input type="checkbox"/> Data Service
<input type="checkbox"/> Views
<input type="checkbox"/> Query and Index Services
<input type="checkbox"/> Search Service
<input type="checkbox"/> Analytics Service

Lab

# Lab 1: Load the Project into NetBeans (optional)

---

Perform the following steps:

- Start a Server X on your Oracle Virtual Box from a Terminal.

`startxfce4`

- Start NetBeans IDE from the Desktop.
- Open the Existing projects “Lab”, “LabSolutions” (solutions) in NetBeans IDE
- Open the “LabSolution” project => This is the solution.
- Open the “Lab” project => This is where you start.

Lab

# Lab 1: Build & Run

---

Update dependency on the libcouchbase:

- Open **java/Lab/pom.xml**
- Change Couchbase Java client dependency to the latest

Build a simple executable jar:

```
cd cb-workshop-2d/java/Lab  
mvn clean compile assembly:single
```

Run main class:

```
java -classpath target/CbDevWorkshop-0.0.1-SNAPSHOT-jar-with-dependencies.jar -  
Dcbworkshop.clusteraddress=<Cluster IP> -Dcbworkshop.user=travel -  
Dcbworkshop.password=couchbase -Dcbworkshop.bucket=travel-sample com.cbworkshop.MainLab
```

Lab

2

# Java SDK Connecting to Couchbase

# Connection Basics

```
// Java
import com.couchbase.client.java.CouchbaseCluster;
import com.couchbase.client.java.CouchbaseBucket;

CouchbaseEnvironment env = DefaultCouchbaseEnvironment.builder()
    .connectTimeout(15000)
    // more custom configuration
    .build();

CouchbaseCluster cluster = CouchbaseCluster.create(env,
    "<host1>", ..., "<hostN>");
cluster.authenticate("<username>", "<password>");
CouchbaseBucket bucket = cluster.openBucket("<bucket-name>");
```

- Reuse cluster and bucket objects, e.g. as singletons
- Make sure to provide more than one hostIP for high availability

# Lab 2: Couchbase Connection

- Use source file: MainLab.java
- Implement method  
initConnection()
- Read config values from System properties:
  - cbworkshop.clusteraddress
  - cbworkshop.user
  - cbworkshop.password
  - cbworkshop.bucket
- Connect to the bucket with the given credentials
- Run application
- Check output:

```
Nov 20, 2017 1:44:25 PM com.couchbase.client.core.CouchbaseCore <init>
INFO: CouchbaseEnvironment: {sslEnabled=false, sslKeystoreFile='null',
sslTruststoreFile='null', sslKeystorePassword=false, sslTruststorePassword=false,
sslKeystore=null, sslTruststore=null, bootstrapHttpEnabled=true,
bootstrapCarrierEnabled=true, bootstrapHttpDirectPort=8091,
bootstrapHttpSslPort=18091, . . .}
```

```
Nov 20, 2017 1:44:33 PM
com.couchbase.client.core.config.DefaultConfigurationProvider$8 call
INFO: Opened bucket travel-sample
```

```
Nov 20, 2017 1:44:35 PM com.couchbase.client.core.node.CouchbaseNode signalConnected
INFO: Connected to Node 127.0.0.1/localhost
```

Lab

3

# Java SDK Key-Value Operations

# Creating Documents

---

- Data can be flat or complex
- Document keys can be custom, automatically generated, or incrementing
- The 'insert' operator will create new documents if the key does not already exist
- The 'upsert' operator will create or replace

```
JsonObject data = JsonObject.create()
    .put("firstname", "Nic")
    .put("lastname", "Raboy");
JsonArray address = JsonArray.create()
    .add(JsonObject.create().put("city", "Mountain View").put("state", "CA"))
    .add(JsonObject.create().put("city", "San Francisco").put("state", "CA"));
data.put("address", address);
bucket.insert(JsonDocument.create("person-1", data));
```

# Retrieving Documents by Key

---

- Data can be retrieved using a key-value lookup or with a N1QL query
- Lookups are significantly faster than indexed queries with N1QL

```
// Java  
bucket.get("person-1").content();
```

# Lab 3: Create Object

- Implement method: `create (String[] words)`
- Compose a JSON document like this:
  - Use the command line parameters from words:
    - document key
    - from
    - to
  - Compose key with prefix "msg::" + provided key
  - Set timestamp to `System.currentTimeMillis()`
  - Set type to "msg"
- Use insert
  - Try several times. See results in console
  - Try same key (Error should appear!)
- Try upsert instead of insert

Key:

**msg::some\_text**

```
{  
  "timestamp": 1511184840248,  
  "from": "luis",  
  "to": "daniel",  
  "type": "msg"  
}
```

Lab

# Lab 4: Read Object

---

- Implement method: `read(String[] words)`
- Use the command line parameter:
  - Document key
- Read the document
- Write the json string to `System.out`
- Test with values:
  - `airline_10226`
  - `route_10009`
  - `hotel_10904`

```
# read airline_10226
{"country": "United States", "iata": "A1", "callsign": "atifly", "name": "Atifly", "icao": "A1F", "id": 10226, "type": "airline"}
```

- **Extra Bonus:** implement code to output a friendly message when document is not found

Lab

# Lab 5: Update Object

---

- Implement method: update (String [ ] words)
- Use the command line parameters:
  - Document key (prefix with “airline\_” in code)
- Read the document
- Modify attribute “name”: set the same value converted toUpperCase
- Use replace to modify

```
# read airline_10642
{"country": "United Kingdom", "iata": null, "callsign": null, "name": "Jc royal.britannica", "icao": "JRB", "id": 10642, "type": "airline"}
# update 10642
# read airline_10642
{"country": "United Kingdom", "iata": null, "callsign": null, "name": "JC ROYAL.BRITANNICA", "icao": "JRB", "id": 10642, "type": "airline"}
```

Lab

# Lab 6: Delete Object

---

- Implement method: delete(String[] words)
- Use the command line parameter:
  - Document key (prefix with “msg::” in code)
- Delete document
- Tip: use create, then delete same key
- Try to read it to test if it is actually deleted

```
# create 1001 luis ana
# read msg::1001
{"from":"luis","to":"ana","type":"msg","timestamp":1511192232720}
# delete 1001
# read msg::1001
java.lang.NullPointerException
#         at com.cbworkshop.MainLab.read(MainLab.java:95)
#         at com.cbworkshop.MainLab.process(MainLab.java:60)
#         at com.cbworkshop.MainLab.main(MainLab.java:30)
```

Lab

4

# Java SDK Subdocument API

# The Goal: Working with Parts of a Document

---

- Get parts of a JSON Document
- Update individual JSON attributes in a document
- Batch subdocument operations together

# Large Documents

---

```
key: nraboy
```

```
{  
  "profile": {  
    "firstname": "Nic",  
    "lastname": "Raboy"  
  },  
  "data": [  
    // 20MB of data  
  ]  
}
```

# Get Part of a Document

---

```
// Java
DocumentFragment<Lookup> result = bucket
    .lookupIn("nraboy")
    .get("profile")
    .execute();
```

# Update Part of a Document

---

```
// Java
SubdocOptionsBuilder builder = new SubdocOptionsBuilder();
builder.createParents(true);
bucket().mutateIn("nraboy")
    .upsert("profile.firstname", "Nicolas", builder)
    .execute();
```

# Chain Subdocument Operations

```
DocumentFragment<Mutation> result = bucket
    .mutateIn("nraboy")
    .replace("profile.firstname", "Nic")
    .insert("profile.gender", "Male", builder)
    .remove("data")
    .withDurability(PersistTo.MASTER, ReplicateTo.NONE)
    .execute();
```

```
DocumentFragment<Lookup> result = bucket
    .lookupIn("myKey")
    .get("sub.value")
    .exists("fruits")
    .execute();
```

```
String subValue = result.content("sub.value", String.class);
boolean fruitsExist = result.content("fruits", Boolean.class);
```

# Lab 7: SubDocument API example

- Implement method: `subdoc(String[] words)`
- Use the command line parameter:
  - Document key (prefix with “msg::” from code)
- Using SubDocument API:
  - Change the actual value of the “from” attribute to “administrator”
  - Add a new attribute: “reviewed”, with value `System.currentTimeMillis()`

```
# create 1005 juan santiago
# read msg::1005
{"from": "juan", "to": "santiago", "type": "msg", "timestamp": 1511196994278}
# subdoc 1005
# read msg::1005
{"reviewed": 1511197006619, "from": "Administrator", "to": "santiago", "type": "msg", "timestamp": 1511196994278}
```

Lab

5

# Java SDK Executing N1QL

# Query Raw String vs. DSL

Raw string query

```
N1qlQueryResult queryResult =  
bucket.query(N1qlQuery.simple("SELECT * FROM `travel-sample` LIMIT 10"));
```

Using the query builder

```
import static com.couchbase.client.java.query.Select.select;  
import static com.couchbase.client.java.query.dsl.Expression.*;  
...  
N1qlQueryResult query =  
bucket.query(select("*").from("`travel-sample`").limit(10));
```

Iterate over the query result

```
for (N1qlQueryRow row : queryResult){  
    System.out.println(row.value().toString());  
}
```

# Query with Parameters

Sample code

```
String sourceairport = ...;
String destinationairport = ...;

String queryStr = "SELECT a.name FROM `travel-sample` r JOIN `travel-sample` a ON
KEYS r.airlineid WHERE r.type=\"route\" AND r.sourceairport=$src AND
r.destinationairport=$dst";

JSONObject params = JSONObject.create()
    .put("src", sourceairport)
    .put("dst", destinationairport);

N1qlQuery query = N1qlQuery.parameterized(queryStr, params,
    N1qlParams.build().adhoc(false));
```

# Query Consistency

---

- **not\_bounded** (fastest)
  - Returns data that is currently indexed and accessible by the index or the view.
- **at\_plus**
  - A query submitted with at\_plus consistency level requires all mutations, up to the moment of the scan\_vector (the logical timestamp passed in with at\_plus), to be processed before the query execution can start.
- **request\_plus**
  - Requires all mutations, up to the moment of the query request, to be processed before the query execution can start.

```
// Java
N1qlQueryResult result =
    bucket.query(
        N1qlQuery.simple(
            statement,
            N1qlParams.build() .consistency(ScanConsistency.AT_PLUS)
        )
    );
```

# Lab 8: Simple Query

- Implement method: query (String[] words)
- Execute the query: "SELECT \* FROM `travel-sample` LIMIT 10"
- Print the results to STDOUT
- Use both implementations, raw and query builder

```
# query
{"travel-sample": {"country": "United States", "iata": "Q5", "callsign": "MILE-AIR", "name": "40-Mile Air", "icao": "MLA", "id": 10, "type": "airline"}}, {"travel-sample": {"country": "United States", "iata": "TQ", "callsign": "TXW", "name": "Texas Wings", "icao": "TXW", "id": 10123, "type": "airline"}}, {"travel-sample": {"country": "United States", "iata": "A1", "callsign": "atifly", "name": "Atifly", "icao": "A1F", "id": 10226, "type": "airline"}}, {"travel-sample": {"country": "United Kingdom", "iata": null, "callsign": null, "name": "JC ROYAL.BRITANNICA", "icao": "JRB", "id": 10642, "type": "airline"}}, {"travel-sample": {"country": "United States", "iata": "ZQ", "callsign": "LOCAIR", "name": "Locair", "icao": "LOC", "id": 10748, "type": "airline"}}, {"travel-sample": {"country": "United States", "iata": "K5", "callsign": "SASQUATCH", "name": "SeaPort Airlines", "icao": "SQH", "id": 10765, "type": "airline"}}, {"travel-sample": {"country": "United States", "iata": "KO", "callsign": "ACE AIR", "name": "Alaska Central Express", "icao": "AER", "id": 109, "type": "airline"}}, {"travel-sample": {"country": "United Kingdom", "iata": "5W", "callsign": "FLYSTAR", "name": "Astraeus", "icao": "AEU", "id": 112, "type": "airline"}}, {"travel-sample": {"country": "France", "iata": "UU", "callsign": "REUNION", "name": "Air Austral", "icao": "REU", "id": 1191, "type": "airline"}}, {"travel-sample": {"country": "France", "iata": "A5", "callsign": "AIRLINAIR", "name": "Airlinair", "icao": "RLA", "id": 1203, "type": "airline"}}
```

Lab

# Lab 9: Query with parameters

---

- Implement method: queryAirports(String[] words)
- Use the command line parameters:
  - sourceairport
  - destinationairport
- Write a query to find airlines (airline names) flying from sourceairport to destinationairport. Use JOIN
- Use a parametrized query
- **TIP:** Highest traffic airport codes: ATL, ORD, LHR, CDG, LAX, DFW, JFK

```
# queryairports JFK LHR
[{"name": "British Airways"},  
 {"name": "Delta Air Lines"},  
 {"name": "American Airlines"},  
 {"name": "US Airways"},  
 {"name": "Virgin Atlantic Airways"},  
 {"name": "Air France"}]
```

Lab

6

# Java SDK

# Asynchronous programming

# Why asynchronous programming?

---

- Synchronous programming is straightforward, e.g. simple loop to create multiple documents

```
for(JsonDocument doc : docs) {  
    bucket.insert(doc);  
}
```

- But difficult to achieve high throughput
  - E.g. if insert takes 1ms, maximum throughput is 1000 op/s
- Multithreading can increase throughput, but creates a lot of overhead
- Asynchronous programming provides an efficient way to achieve high throughput

# Observable Pattern

- Observable = a stream of data
- RxJava operators to manipulate the stream

	<i>Single</i>	<i>Multiple</i>
<i>Sync ("pull")</i>	<i>T</i>	<i>Iterable&lt;T&gt;</i>
<i>Async ("push")</i>	<i>Future&lt;T&gt;</i>	<i>Observable&lt;T&gt;</i> 

# Batching with RxJava

---

- Implicit batching is performed by utilizing a few operators:
  - **Observable.just()** or **Observable.from()** to generate an Observable that contains the data you want to batch on.
- **flatMap()** to process the stream events with the Couchbase Java SDK and merge the results asynchronously.
- **last()** to wait until the last event of the stream is received
- **toList()** to transform the events into a single list of results. Useful for reading data
- **toBlocking().single()** transforms the stream into a synchronous call returning the result

# Batching with RxJava

- The following example creates an observable stream of 5 keys to load in a batch,
- asynchronously fires off get() requests against the SDK,
- waits until the last result has arrived,
- and then converts the result into a list and blocks at the very end

```
Cluster cluster = CouchbaseCluster.create();
Bucket bucket = cluster.openBucket();

List<JsonDocument> foundDocs = Observable
    .just("key1", "key2", "key3", "key4", "key5")
    .flatMap(new Func1<String, Observable<JsonDocument>>() {
        @Override
        public Observable<JsonDocument> call(String id) {
            return bucket.async().get(id);
        }
    })
    .toList()
    .toBlocking()
    .single();
```

# Batching with RxJava

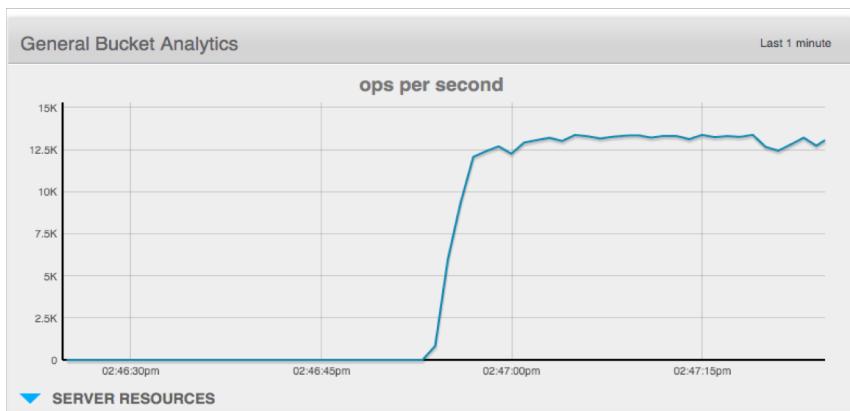
- If you wrap the code in a helper method, you can provide very nice encapsulated batching semantics

```
public List<JsonDocument> bulkGet(final Collection<String> ids) {  
    return Observable  
        .from(ids)  
        .flatMap(new Func1<String, Observable<JsonDocument>>() {  
            @Override  
            public Observable<JsonDocument> call(String id) {  
                return bucket.async().get(id);  
            }  
        })  
        .toList()  
        .toBlocking()  
        .single();  
}
```

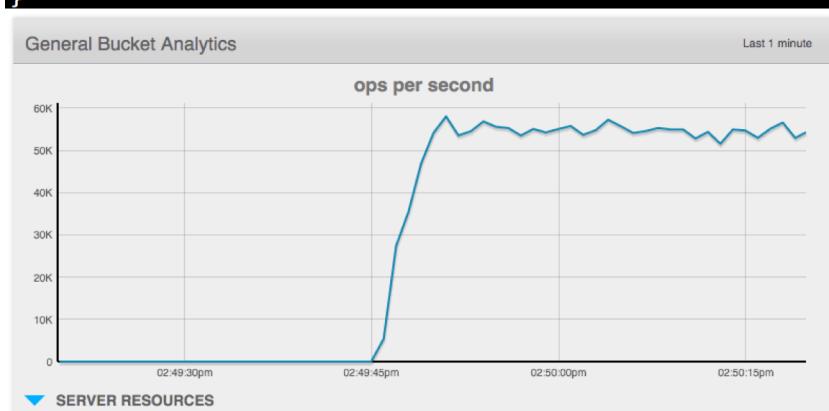
# Batching with RxJava

Here are two code samples, both synchronous, that showcase serialized and batched loading of documents.

```
// Serialized workload of loading documents
while(true) {
    List<JsonDocument> loaded = new ArrayList<JsonDocument>();
    int docsToLoad = 10;
    for (int i = 0; i < docsToLoad; i++) {
        JsonDocument doc = bucket.get("doc--" + i);
        if (doc != null) {
            loaded.add(doc);
        }
    }
}
```



```
// Same workload, utilizing batching effects
while(true) {
    int docsToLoad = 10;
    Observable
        .range(0, docsToLoad)
        .flatMap(new Func1<Integer, Observable<JsonDocument>>() {
            @Override
            public Observable<JsonDocument> call(Integer i) {
                return bucket.async().get("doc--"+i);
            }
        })
        .toList()
        .toBlocking()
        .single();
}
```



# Batching with RxJava - Batching mutations

- The following code generates a number of fake documents and inserts them in one batch.
- Note that you can decide to either collect the results with **toList()** as shown before or just use **last()** as shown here to wait until the last document is properly inserted.

```
// Generate a number of dummy JSON documents
int docsToCreate = 100;
List<JsonDocument> documents = new ArrayList<JsonDocument>();
for (int i = 0; i < docsToCreate; i++) {
    JsonObject content = JsonObject.create()
        .put("counter", i)
        .put("name", "Foo Bar");
    documents.add(JsonDocument.create("doc-"+i, content));
}

// Insert them in one batch, waiting until the last one is done.
Observable
    .from(documents)
    .flatMap(new Func1<JsonDocument, Observable<JsonDocument>>() {
        @Override
        public Observable<JsonDocument> call(final JsonDocument docToInsert) {
            return bucket.async().insert(docToInsert);
        }
    })
    .last()
    .toBlocking()
    .single();
```

# Bulk Read Async

---

```
List<String> priceKeys;
List<JsonObject> res = Observable
    .from(priceKeys)
    .flatMap(k -> bucket.async().get(k))
    .map(doc -> doc.content())
    .toList()
    .toBlocking()
    .single();
```

# Bulk Write Async

---

Sync version:

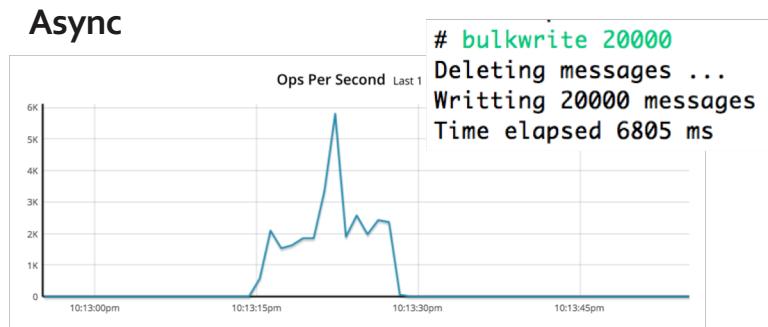
```
for(JsonDocument doc : docs) {  
    bucket.insert(doc);  
}
```

Async version:

```
List<JsonDocument> docs = ...;  
Observable  
    .from(docs)  
    .flatMap(doc -> bucket.async().insert(doc))  
    .last()  
    .toBlocking()  
    .single();
```

# Lab 10: Bulk Write Performance

- Implement method: bulkWrite (String[] words) : Async version
- Implement method: bulkWriteSync (String[] words) : Sync version
- Read parameters from command line:
  - size: Number of messages to insert. Keys will be from msg::1 to msg::[size]
- Delete all messages in the bucket: DELETE FROM `travel-sample` WHERE type="msg"
- Create an array of JsonDocuments of messages
- Insert the messages into the bucket (in async / sync way)
- Print the time elapsed to STDOUT
- Compare results sync vs. async. Check both time and operations per second in the console



Lab

# Query Async mode

```
bucket.async()
    .query(select("*").from(i(`travel-sample`)).limit(5))
    .flatMap(result ->
        result.errors()
        .flatMap(e ->
            Observable.<AsyncN1qlQueryRow>error(
                new CouchbaseException("N1QL Error/Warning: " + e)))
        .switchIfEmpty(result.rows())
    )
    .map(AsyncN1qlQueryRow::value)
    .subscribe(
        rowContent -> System.out.println(rowContent),
        runtimeError -> runtimeError.printStackTrace()
    );
}
```

- Unlike the synchronous method, does not block the calling thread
- The query results are processed asynchronously as they arrive by the subscribe handlers

# Lab 11: Simple Query – Async version

---

- Implement method: queryAsync (String[] words)
- Execute the query: “SELECT \* FROM `travel-sample` LIMIT 5”
- Print the results to STDOUT
- Use asynchronous implementation

```
# queryasync
# {"travel-sample": {"country": "United States", "iata": "Q5", "callsign": "MILE-AIR", "name": "40-Mile Air", "icao": "MLA", "id": 10, "type": "airline"}}
# {"travel-sample": {"country": "United States", "iata": "TQ", "callsign": "TXW", "name": "Texas Wings", "icao": "TXW", "id": 10123, "type": "airline"}}
# {"travel-sample": {"country": "United States", "iata": "A1", "callsign": "atifly", "name": "Atifly", "icao": "A1F", "id": 10226, "type": "airline"}}
# {"travel-sample": {"country": "United Kingdom", "iata": null, "callsign": null, "name": "JC ROYAL.BRITANNICA", "icao": "JRB", "id": 10642, "type": "airline"}}
# {"travel-sample": {"country": "United States", "iata": "ZQ", "callsign": "LOCAIR", "name": "Locair", "icao": "LOC", "id": 10748, "type": "airline"}}
```

Lab

# Thank you



Couchbase